

Relational Table design from EER diagrams

(Mapping from EER to Relational Model)

- Goal of table design
- Principles for table design
 - 1. From information capture pt. of view**
 - 2. From systems (space,time) and programming effort point of view**
- Rules for table design start from 1 and then are refined using 2.
- Some worked examples

Minimal needs for declaring tables:

- list of columns (and their datatype)
- primary key
- foreign key references
- non-null constraints
- other constraints

A simple way to start:

Create a separate table for every entity and relationship appearing in the EER diagram.

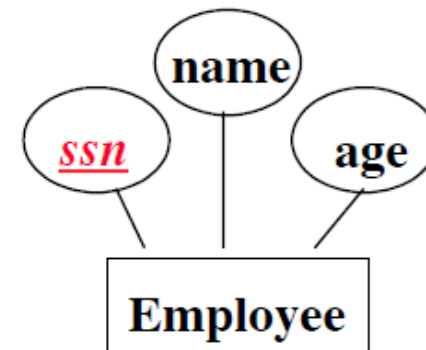
We will go through the various kinds of entities and relationships, showing an example design

(BUT we will have to come back and revisit some of these, in light of other principles of table design, to be discussed a bit later)

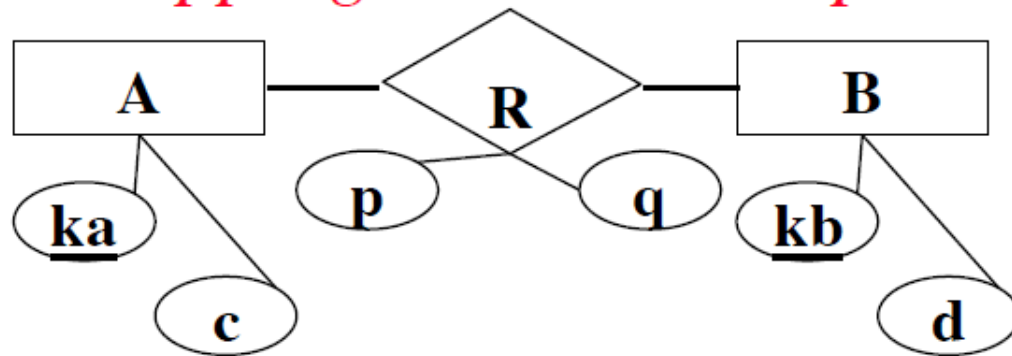
Mapping an Entity Set:

- Relation T_A for every entity A. Attributes of relation are all attributes of the entity.
- Constraints
 - » **key of ER entity becomes primary key**

```
table T_Employee(  
    ssn integer primary key,  
    name varchar(30),  
    age integer  
)
```



Mapping a Relationship:



e.g., A=Person, B=Car
R=buys

table T_R

(ka integer *references* T_A

kb char(5) *references* T_B

p date,

q char(3),

primary key (ka , kb))

on delete cascade,
on delete cascade,

- same idea for n-ary relationships, which have n foreign keys then

Mapping a Relationship e.g.:

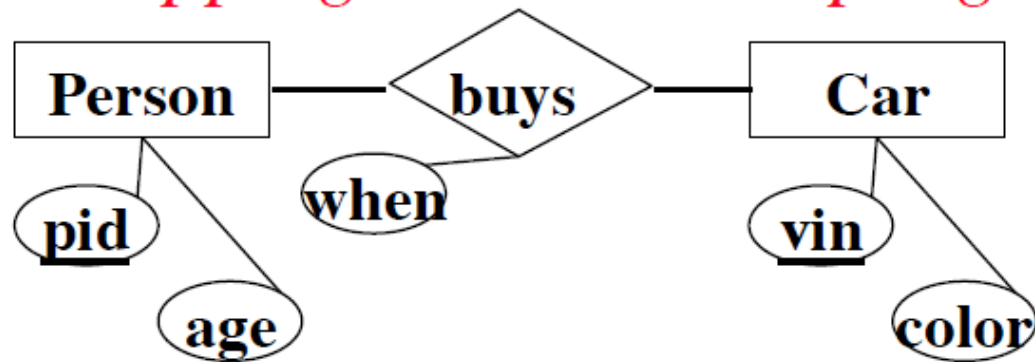


table T_buys

(pid integer *references* T_Person

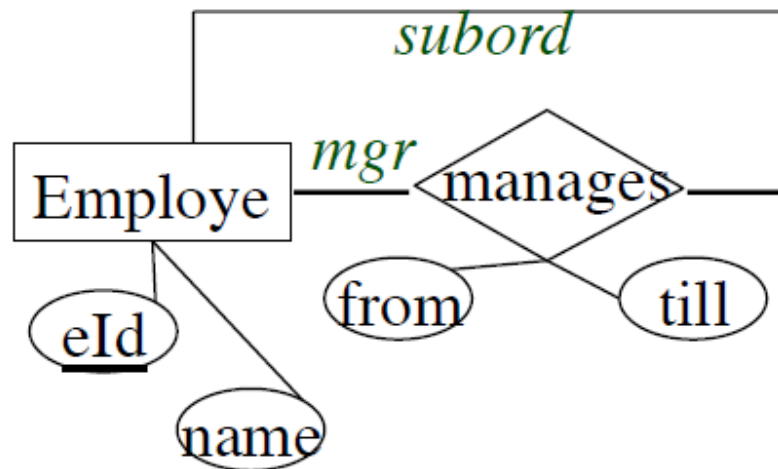
vin char(25) *references* T_Car

when date,

primary key (pid, vin))

on delete cascade,
on delete cascade,

Mapping a “reflexive” relationship (when $A = B$)



`T_Employee(eId, name)`

`T_manages(mgr_eId, subord_eId, from, till)`

Mapping Relationships with functional constraints (arrow):

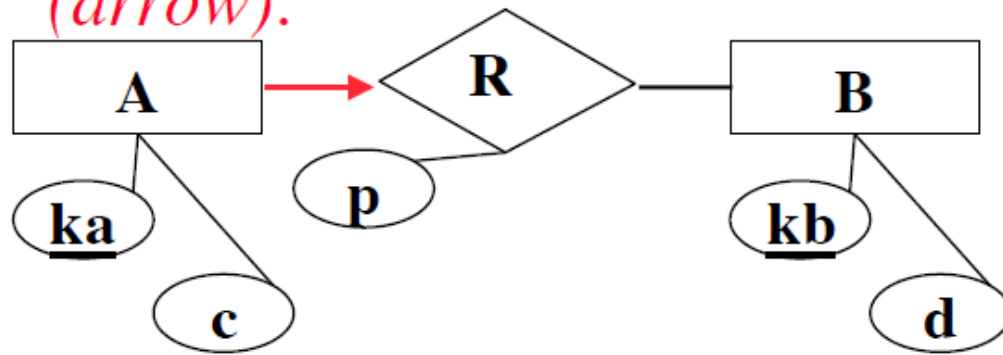


table T_R

(ka char(25) *references* T_A,

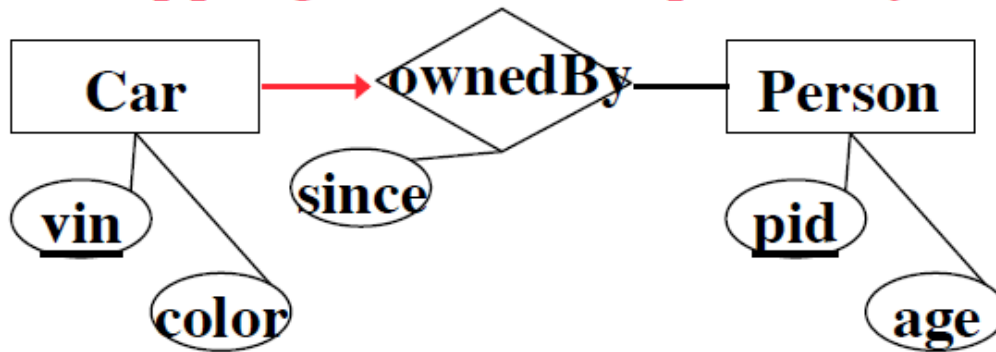
kb int *references* T_B,

p date,

primary key (ka))

- The key had to be made smaller, since for each A object there is at most one R-related B object, so (ka, kb) is a super-key but not a key.
- Note: this way the functional constraint is checked by the schema!

Mapping Relationships with functional constraints (eg):



```
table T_ownedBy
  ( vin char(25) references T_Car,
    pid int references T_Person,
    since date,
    primary key ( vin ) )
```

Mapping 1-N Relationships vs N:M recap

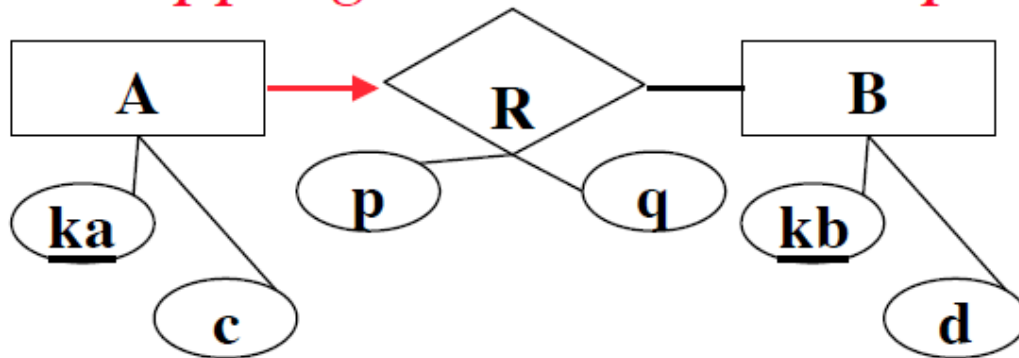


table T_R(ka,kb,p,q)
has *primary key* (ka)

In general the key of the entity with the “arrow” to the binary relationship is the key to the relationship table itself.

(Suppose B or A's participation in R is 'total' (thick line)
Note that there is no constraint on table T_R alone that
can enforce this.)

Mapping 1-1 Relationships:

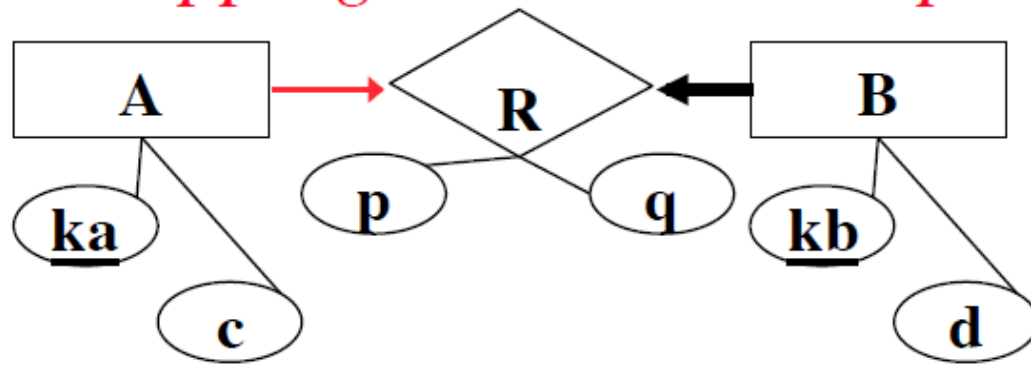
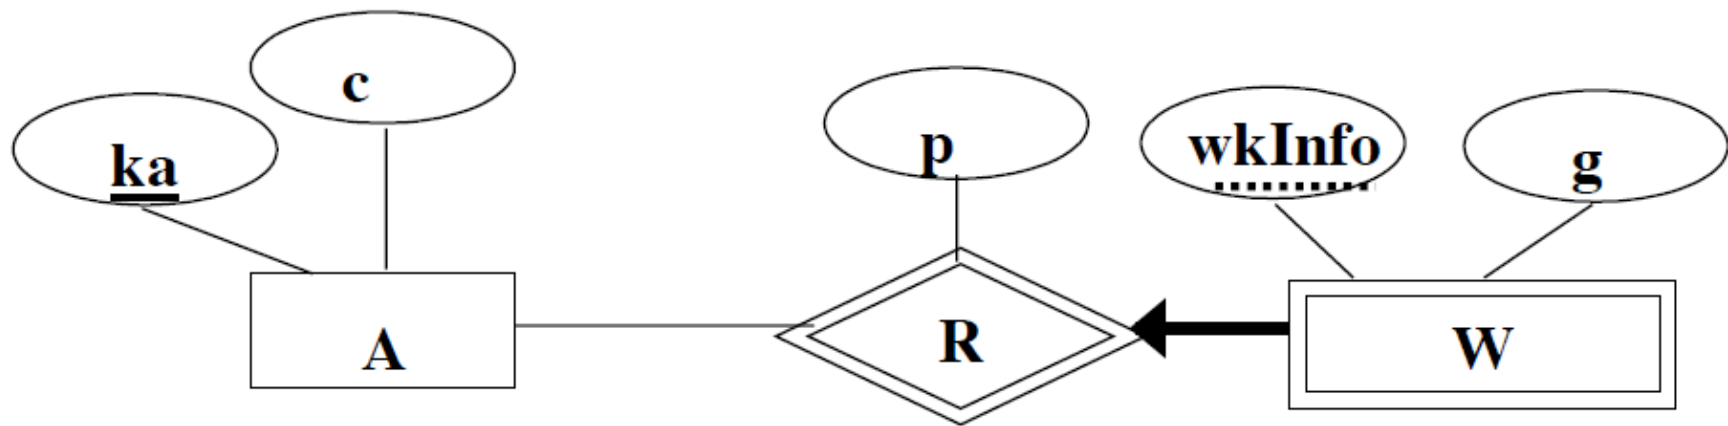


table $T_R(\underline{ka}, kb, p, q)$ or $T_R(ka, \underline{kb}, p, q)$
are acceptable alternatives.

If one entity's participation is total (say B), it is better to choose that one to be the primary key (kb). (See later how NOT NULL can then enforce this constraint in a merged table.)

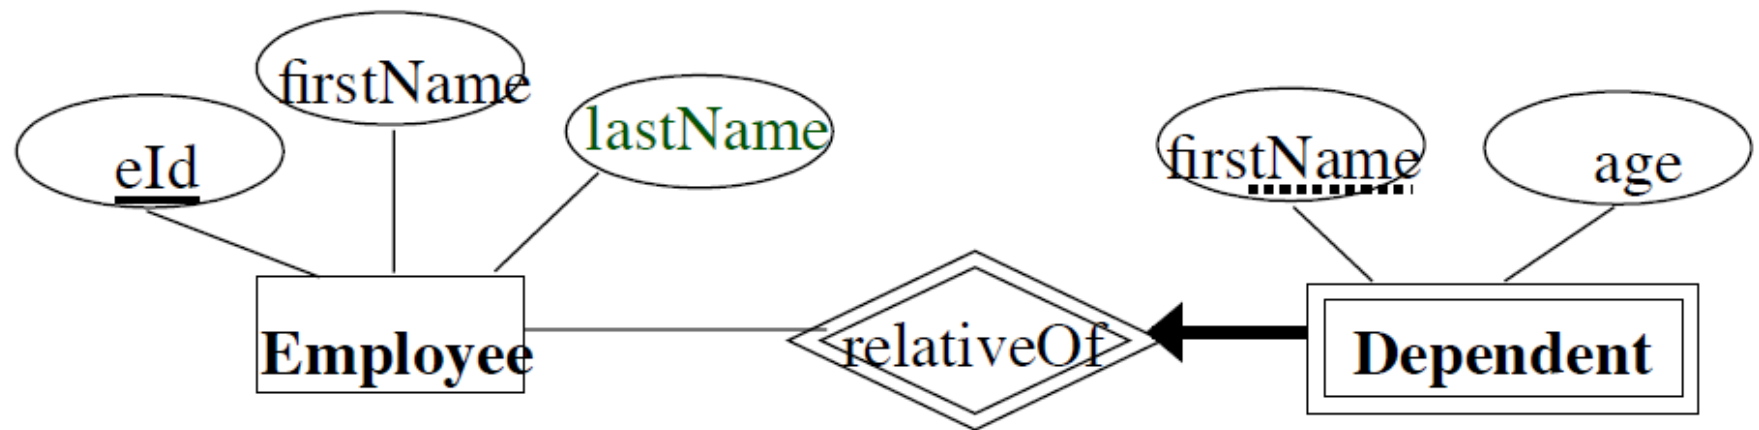
Mapping Weak Entities



```
table T_W(  
  wkInfo  
  g  
  p  
  ka references T_A  
  primary key (wkInfo, ka)
```

- Cannot have *wkInfo* be a key alone since it does not identify objects in *W* on its own. (That is the whole point of weak entities)
- Throw in attribute *p* of the relationship, so *T_W* is really *T_WR*

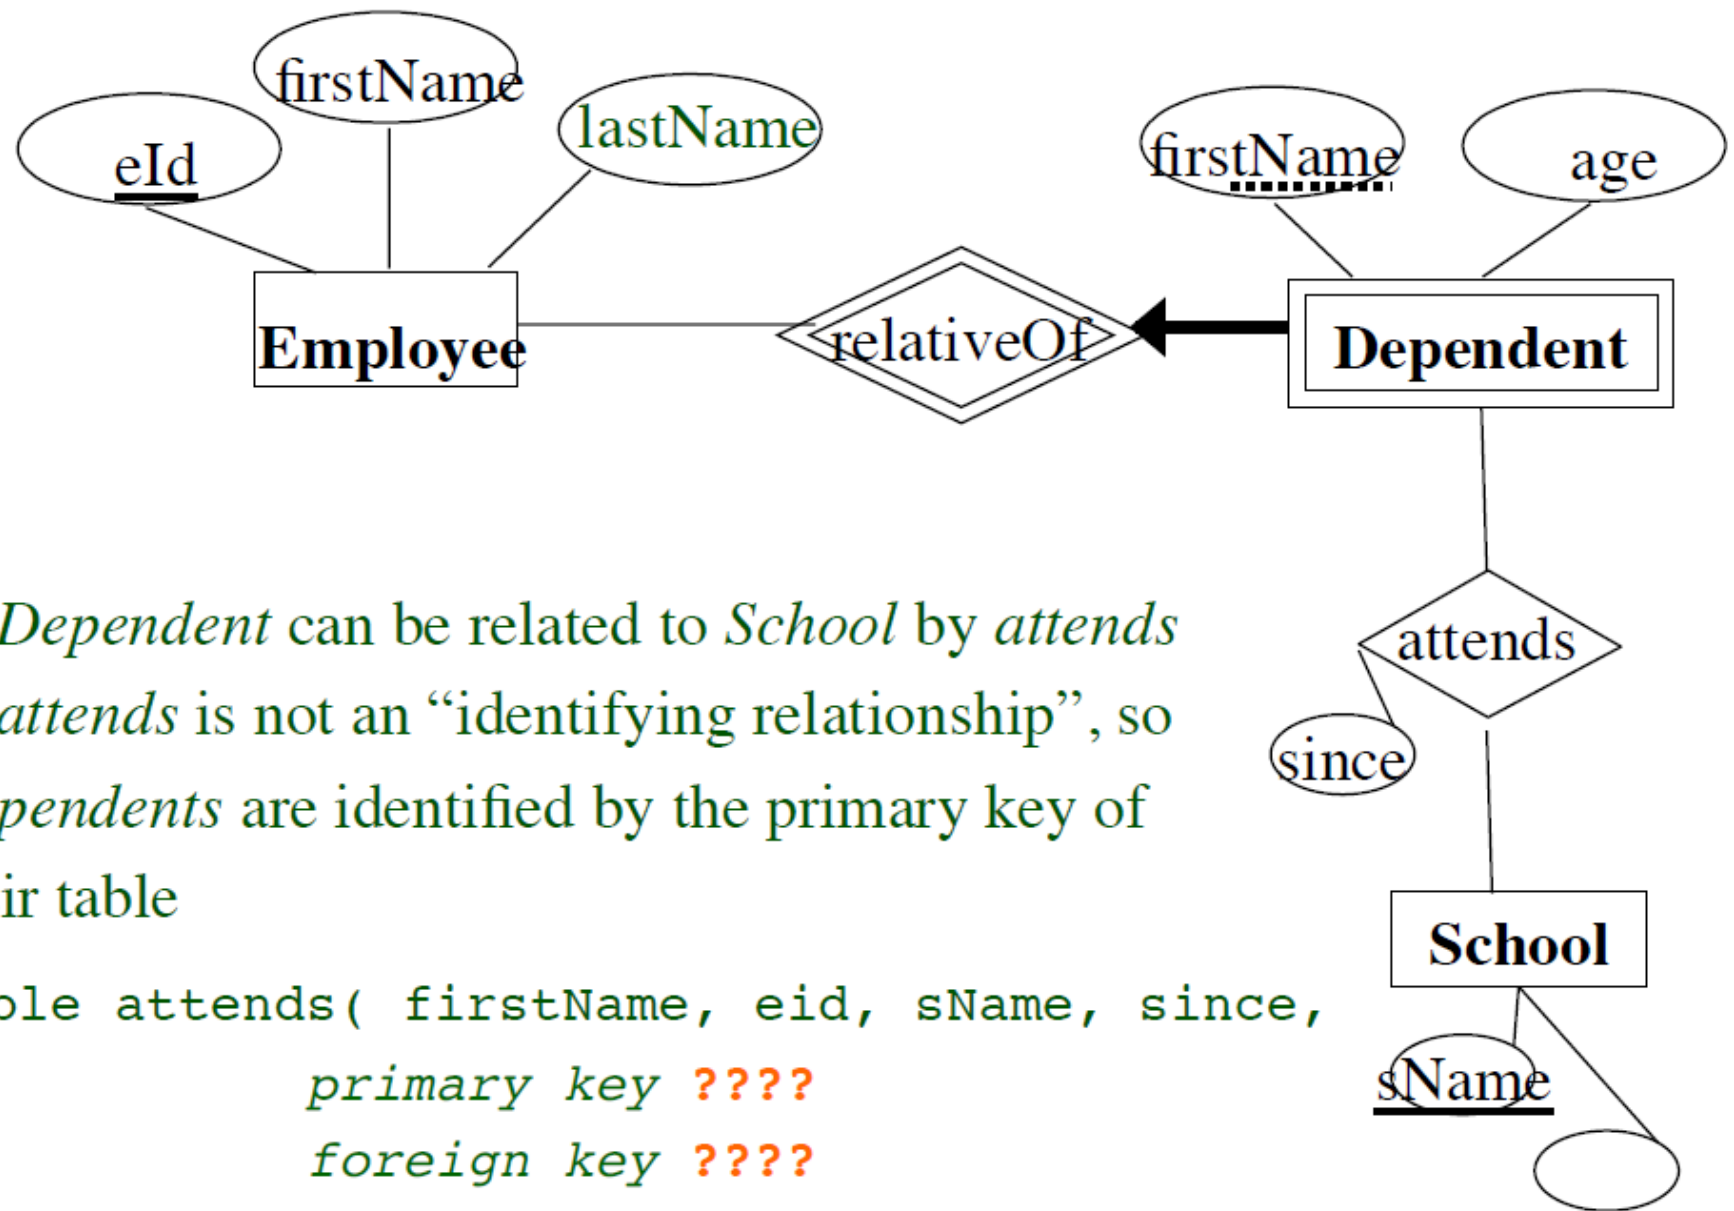
Weak entity example:



```
table T_Dependent(  
  firstName,  
  eId,  
  foreign key (eId) references T_Employee,  
  primary key (firstName, eId) )
```

Describes a Dependent

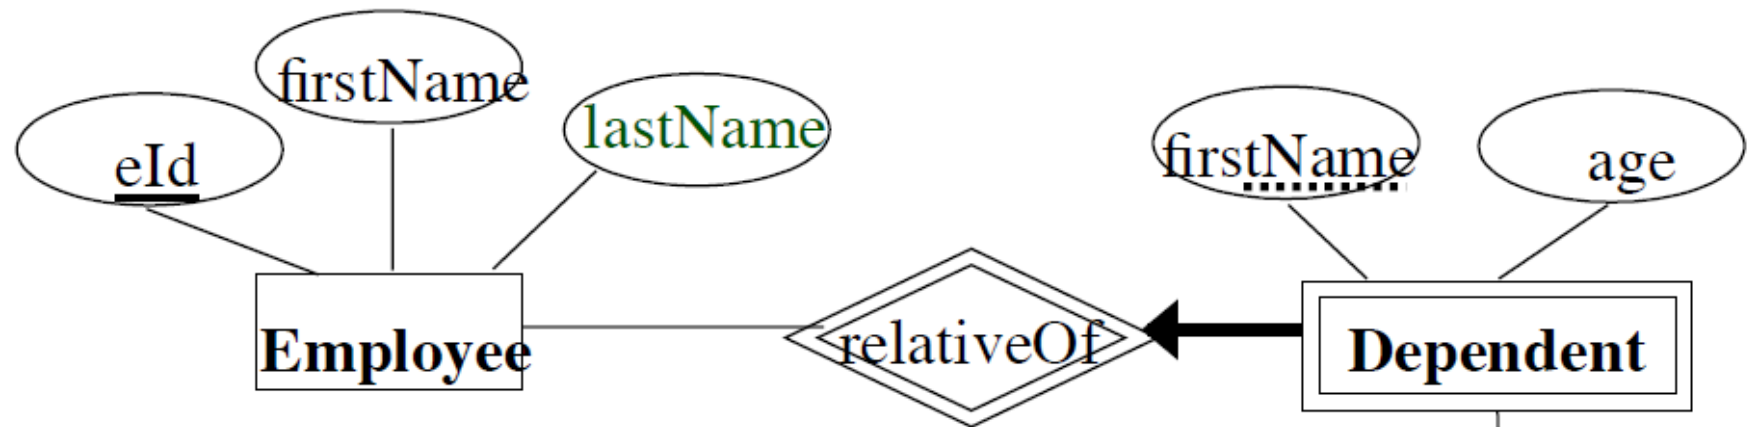
Weak entities participating in other relationships:



- *Dependent* can be related to *School* by *attends*
- *attends* is not an “identifying relationship”, so *Dependents* are identified by the primary key of their table

```
table attends( firstName, eid, sName, since,  
              primary key ????,  
              foreign key ????)
```

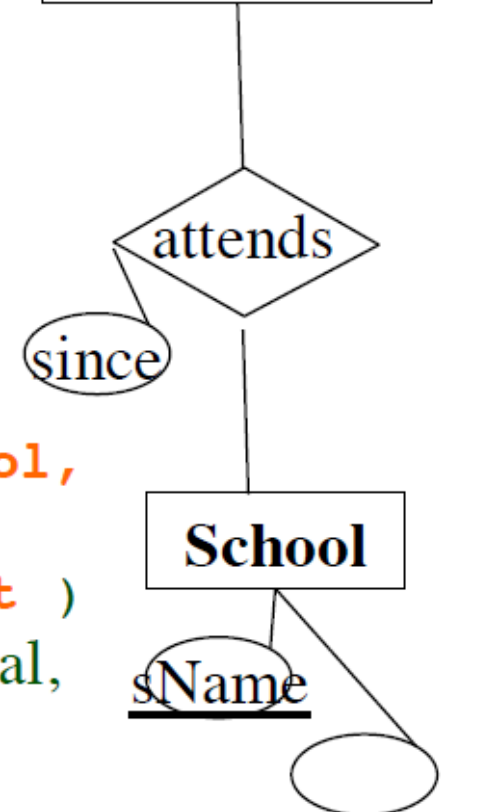
Weak entities participating in other relationships:



```

table attends(
  firstName,
  eid,
  sName,
  since,
  primary key (firstName, eid, sName),
  foreign key (sname) references School,
  foreign key (firstName, eid) refs
                                Dependent )
  
```

Note that “firstName references Dependent” is illegal,
and “eid reference Employee” is too weak.

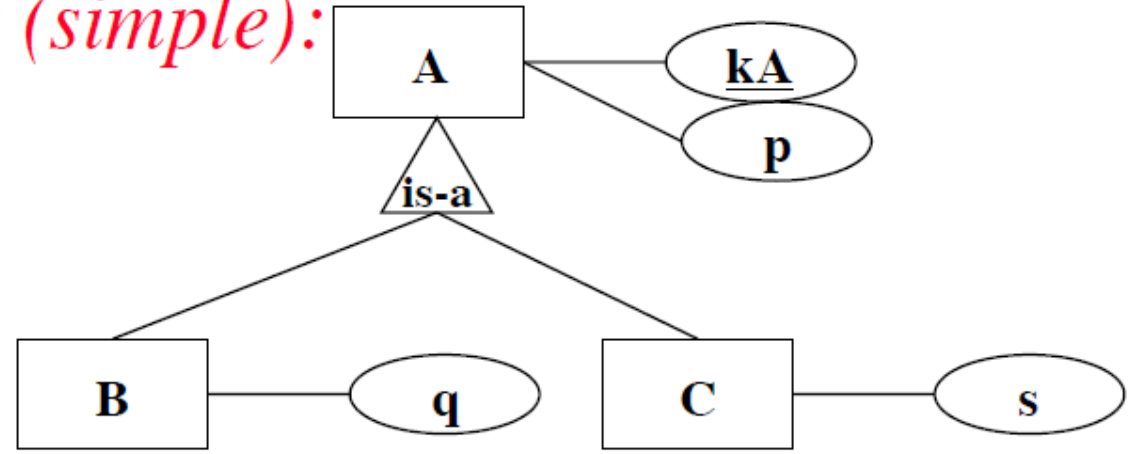


Mapping Subclasses (simple):

```
table T_A(  
    kA primary key  
    p )
```

```
table T_B(  
    kA primary key  
    q)
```

```
table T_C(  
    kA primary key  
    s)
```



foreign key(kA) references T_A

foreign key(kA) references T_A

These capture the
subclass
inclusion constraints
“every B,C is an A”

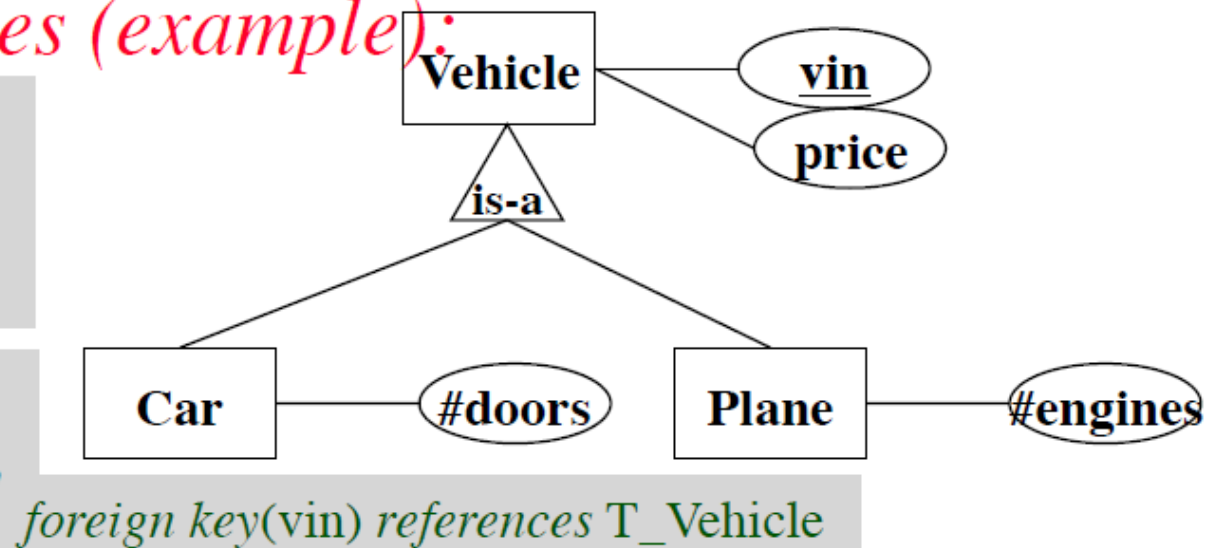
- One table per entity, just like before
- For B and C, just inherit the same key as you did for the root class (A here), since every B and C is also an A
- Disjointness or covering by subclasses need to be stated as *check* constraints *Exercise: try writing this!!*

Mapping Subclasses (example):

```
table T_Vehicle(  
    vin primary key  
    price)
```

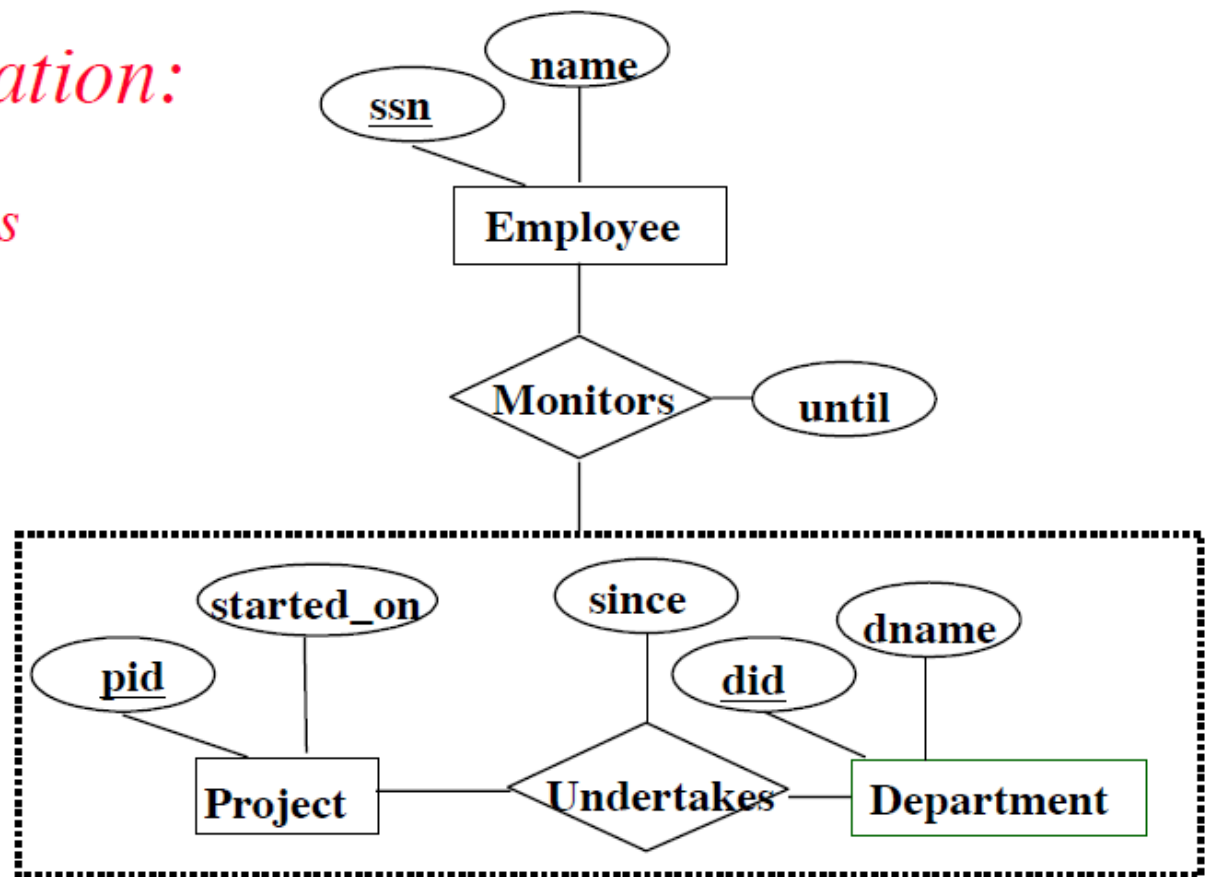
```
table T_Car(  
    vin primary key  
    #doors)
```

```
table T_Plane(  
    vin primary key  
    #engines)
```



- One table per entity, just like before
- For B and C, just inherit the same key as you did for the root class (A here), since every B and C is also an A
- Disjointness or covering by subclasses need to be stated as *check* constraints *Excercise: try writing this!!*

*Tables for Aggregation:
nothing new:
just apply previous ideas*



Tables for Aggregation: just apply previous ideas!

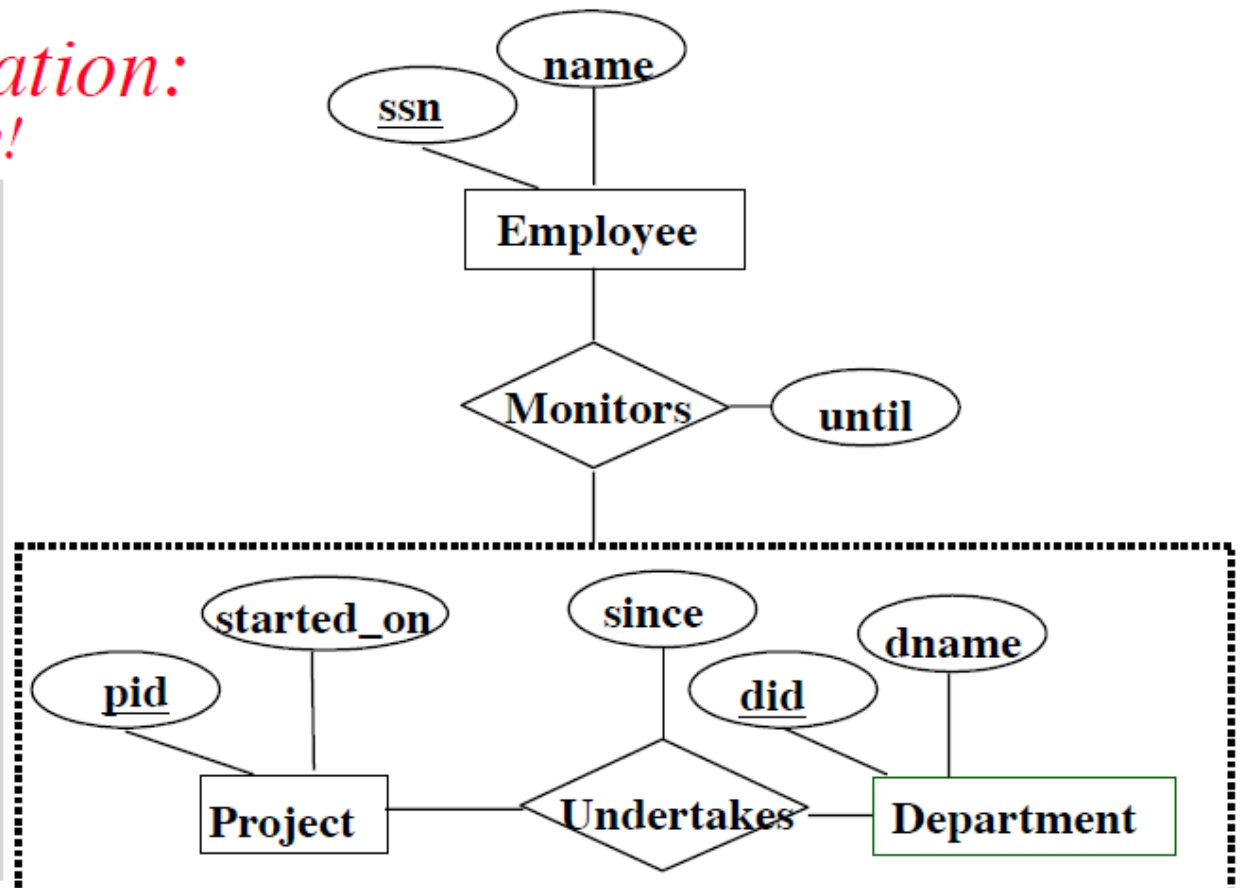
Tables for nested reln's

T_proj(pid, started_On)

T_dept(did, dname)

T_undertakes(
pid,did, since)

T_employee(ssn, name)



- Table for “aggregate” relationship treats Undertakes as entity

```
table T_monitors(  
  ssn integer references T_emp  
  did integer  
  pid integer  
  foreign key (did,pid) references T_undertakes  
  until date  
  primary key (ssn,did,pid) )
```

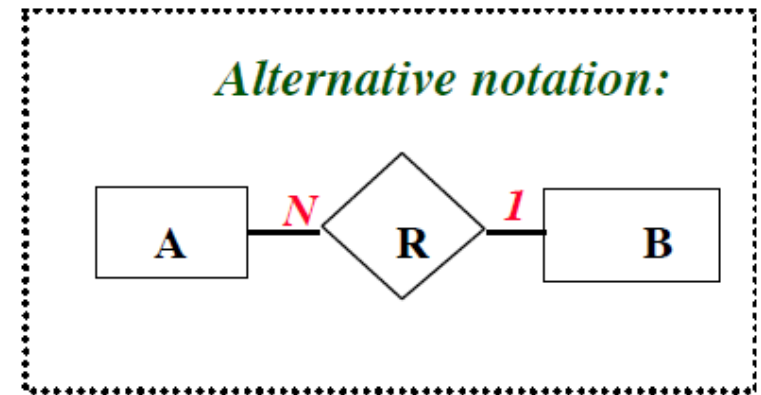
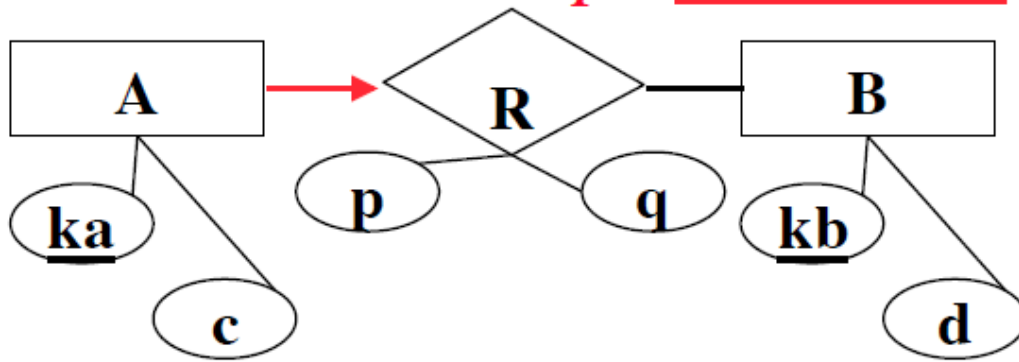
DRAW >>

Merge rule:

If you have two tables of the form $T(\underline{K}, X)$ and $S(\underline{K2}, Y)$, Y not null, where $K2$ is a foreign key referencing $T(K)$ then you can merge S into T to get instead a single table $TS(\underline{K}, X, Y)$. Column Y will have NULLs for all keys in T but not S .

- We'll talk later why this is good
- The reason we need Y to be not null, is to be able to distinguish the case when a tuple was in T only, but not in S . (Y will be null in table TS in exactly this case.)
- If no such attribute Y is present, a boolean attribute $inS?$ can be introduced to distinguish the tuples in TS which were not in S .
- One can think of TS as being the “left outer join” of T and S (the plain join would lose the tuples in T that are not in S)

1-N Relationships Revisited



- **Merge** $T_R(\underline{ka}, kb, p, q)$ **into** $T_A(\underline{ka}, c)$

```

table T_A
    ( ka
      c
      R_kb references T_B
      p
      q
      primary key ( ka )
    )
  
```

Describes an A

*Describes the
B related to A
via R*