


Chapter 3

Transport Layer



Computer Networking: A Top Down Approach

7th edition

Jim Kurose, Keith Ross

Pearson/Addison Wesley


April 2016

Chapter 3: Transport Layer

our goals:

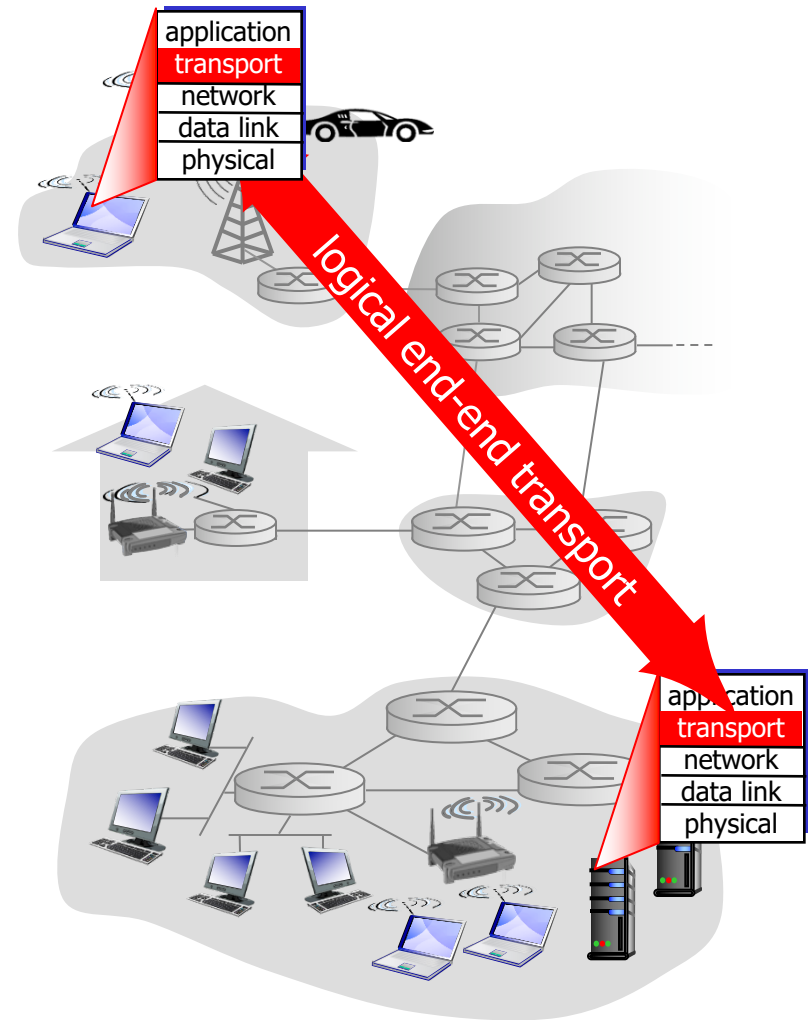
- understand principles behind transport layer services:
 - multiplexing, demultiplexing
 - flow control
 - congestion control
- learn about Internet transport layer protocols:
 - UDP: connectionless transport
 - TCP: connection-oriented reliable transport

Internet Protocol Stack



Transport services and protocols

- provide *logical communication* between **app processes** running on **different hosts**
- transport protocols run in end systems
 - **send side:** breaks app messages into *segments*, passes to network layer
 - **rcv side:** reassembles segments into messages, passes to app layer
- more than one transport protocol available to apps
 - Internet: **TCP and UDP**



Internet transport-layer protocols

■ Transmission Control Protocol (**TCP**)

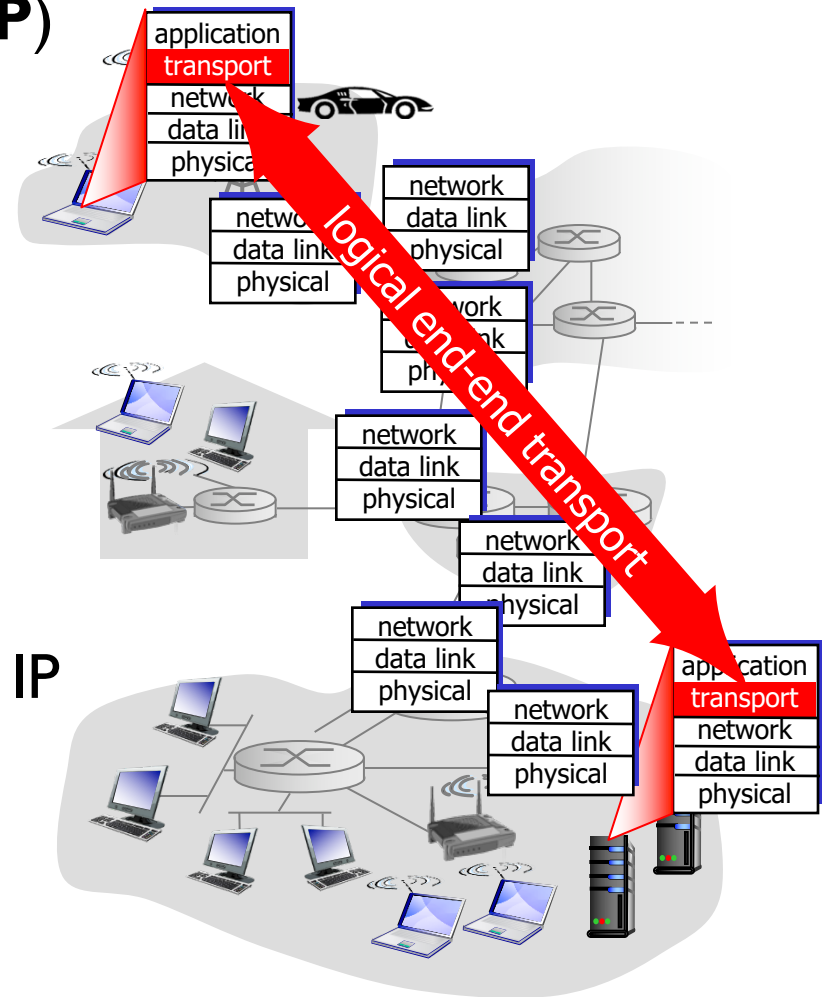
- reliable, in-order delivery
- congestion control
- flow control
- connection setup

■ User Datagram Protocol: **UDP**

- unreliable, unordered delivery
- Simple extension of “best-effort” IP

■ Services not available:

- delay guarantees
- bandwidth guarantees



Multiplexing/demultiplexing

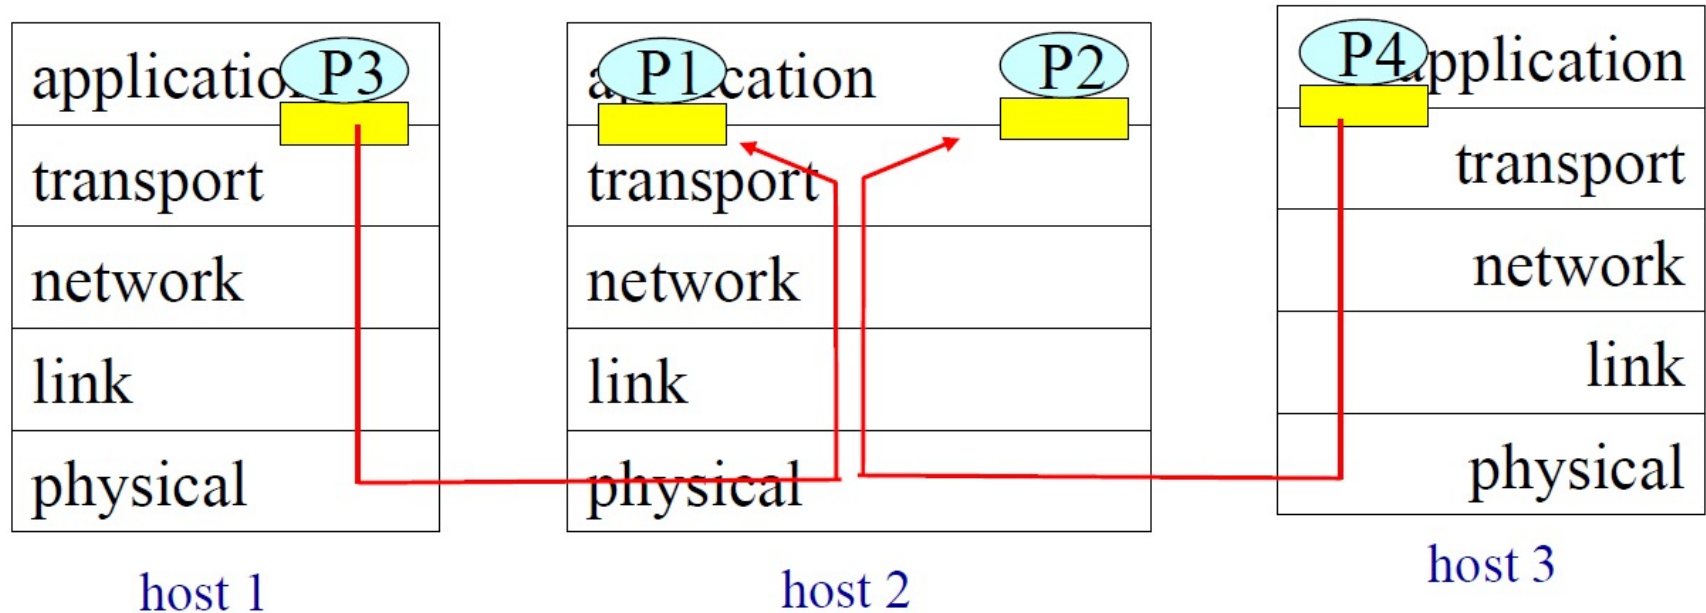
Multiplexing at send host:

gathering data from multiple sockets, enveloping data with header (later used for demultiplexing)

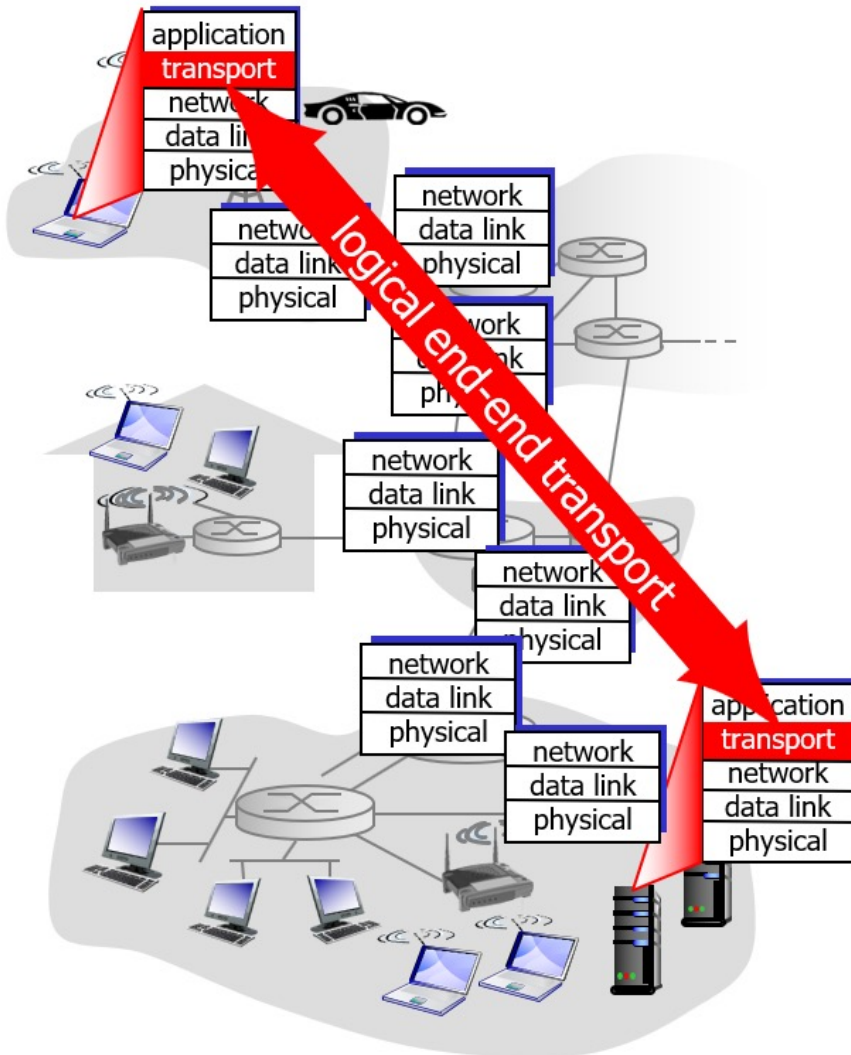
Demultiplexing at rcv host:

delivering received segments to correct socket

 = socket  = process




UDP: User Datagram Protocol [RFC 768]



- UDP use:
 - streaming multimedia apps (loss tolerant, rate sensitive)
 - DNS
- reliable transfer over UDP:
 - add reliability at application layer
 - application-specific error recovery!

UDP: segment header



UDP segment format

why is there a UDP?

- no connection establishment (which can add delay)
- simple: no connection state at sender, receiver
- small header size
- no congestion control: UDP can blast away as fast as desired

Reliable Data Transfer

- ❑ Problem: Reliability
 - ❖ A reliable link **even though** packets can be corrupted or get lost
- ❑ Where can packets **be corrupted or lost?**
 - ❖ In the network
 - ❖ At the receiver
- ❑ Solution: keep track of the packets
 - ❖ Not as simple as one would expect
 - ❖ Reliable Transmission & Flow Control

Reliable Transmission & Flow Control

- What to do when there is a packet loss?
 - ❖ On the link: Retransmit
 - ❖ At the receiver: Flow Control

- Sender needs to know if a packet was lost, How?
 - ❖ ACK (Positive & Negative)

- Sender needs to retransmit, How?
 - ❖ Timeouts

Some Flow Control Algorithms

1. Flow control for the ideal network
2. Stop and Wait for **noiseless** channels
3. Stop and Wait for **noisy** channels
4. Sliding window protocols
5. Sliding window with **error control**
 - Go Back N
 - Selective Repeat

1. Flow control in the ideal network

Assumptions:


- (1) Error free transmission link,
- (2) Infinite buffer at the receiver

No acknowledgement necessary


Since the data link is **error-free** and the receiver can buffer **as many packet as it likes**, no packet will ever be lost

2. Stop-and-Wait Noiseless Channel


Packet Length = L ; Bandwidth = R ; $RTT = 2 \times \text{Prop Delay}$



3. Stop-and-Wait Noisy Channel



Packet retransmitted




Packet retransmitted

Is Stop and Wait the best we can do?


Stop and Wait is an **effective** form of flow control, but...
It's not very **efficient**.

1. Only **one data frame** can be in transit on the link at a time
2. When waiting for an acknowledgement, the sender **cannot transmit any frames**

Better solution? Pipelined Protocol : Sliding Window




(a) a stop-and-wait protocol in operation




(b) a pipelined protocol in operation

Sliding Window example

Sender window



Receiver window




(a)

(b)

(c)

(d)



- (a) Initial state, no frames transmitted, receiver expects frame 0
 (b) Sender transmits frame 0, receiver buffers frame 0
 (c) Receiver ACKS frame 0
 (d) Sender receives ACK, removes frame 0


This protocol **behaves identically** to stop and wait for a noisy channel

Sliding Window with Maximum Sender Window Size SWS

Sender Window size: The maximum number of frames **the sender may transmit** without receiving any acknowledgements


With a maximum window size of *SWS*, the sender **can transmit up to *SWS* frames** before "being blocked"

This allows the sender to transmit several frames **before** waiting for an acknowledgement




$$LFS - LAR \leq SWS$$


Sender-Side Window with $W_s=2$




(a)




(b)




(c)




(d)




(e)



(f)



(g)



(h)

(a) Initial window state

(b) Send frame 0

(c) Send frame 1

(d) ACK for frame 0 arrives

(e) Send frame 2

(f) ACK for frame 1 arrives

(g) ACK for frame 2 arrives, send frame 3


(h) ACK for frame 3 arrives

Sliding Window with Maximum Receiver Window Size


Receiver Window size: The maximum number of frames **the receiver may receive** before returning an acknowledgement to the sender

With a maximum window size of RWS , **the receiver rejects packets** if $SeqNum \leq LFR$ or $SeqNum > LAF$


Why? Outside the window




Receiver-Side Window with $W_R=2$




(a)




(b)




(c)




(d)




(e)



(f)



(g)



(h)

(a) Initial window state

(b) Nothing happens

(c) Frame 0 arrives, ACK frame 0

(d) Nothing happens


(e) Frame 1 arrives, ACK frame 1

(f) Frame 2 arrives, ACK frame 2

(g) Nothing happens

(h) Frame 3 arrives, ACK frame 3

Sender-Side Window with $W_s=2$



Receiver-Side Window with $W_R=2$