


3 Flow Control

TCP Flow Control

- ❑ TCP uses a modified version of the sliding window
- ❑ In acknowledgements, TCP uses the "Window size" field to **tell the sender** how many bytes it may transmit
- ❑ TCP uses **bytes**, not packets, as sequence numbers

TCP Flow Control (cont'd)

Important information in TCP/IP packet headers



Number of
bytes in packet (N)

Sequence number
of first data byte in
packet (SEQ)



ACK bit
set

Sequence number
of next expected
byte (ACK)

Window size
at the receiver
(WIN)

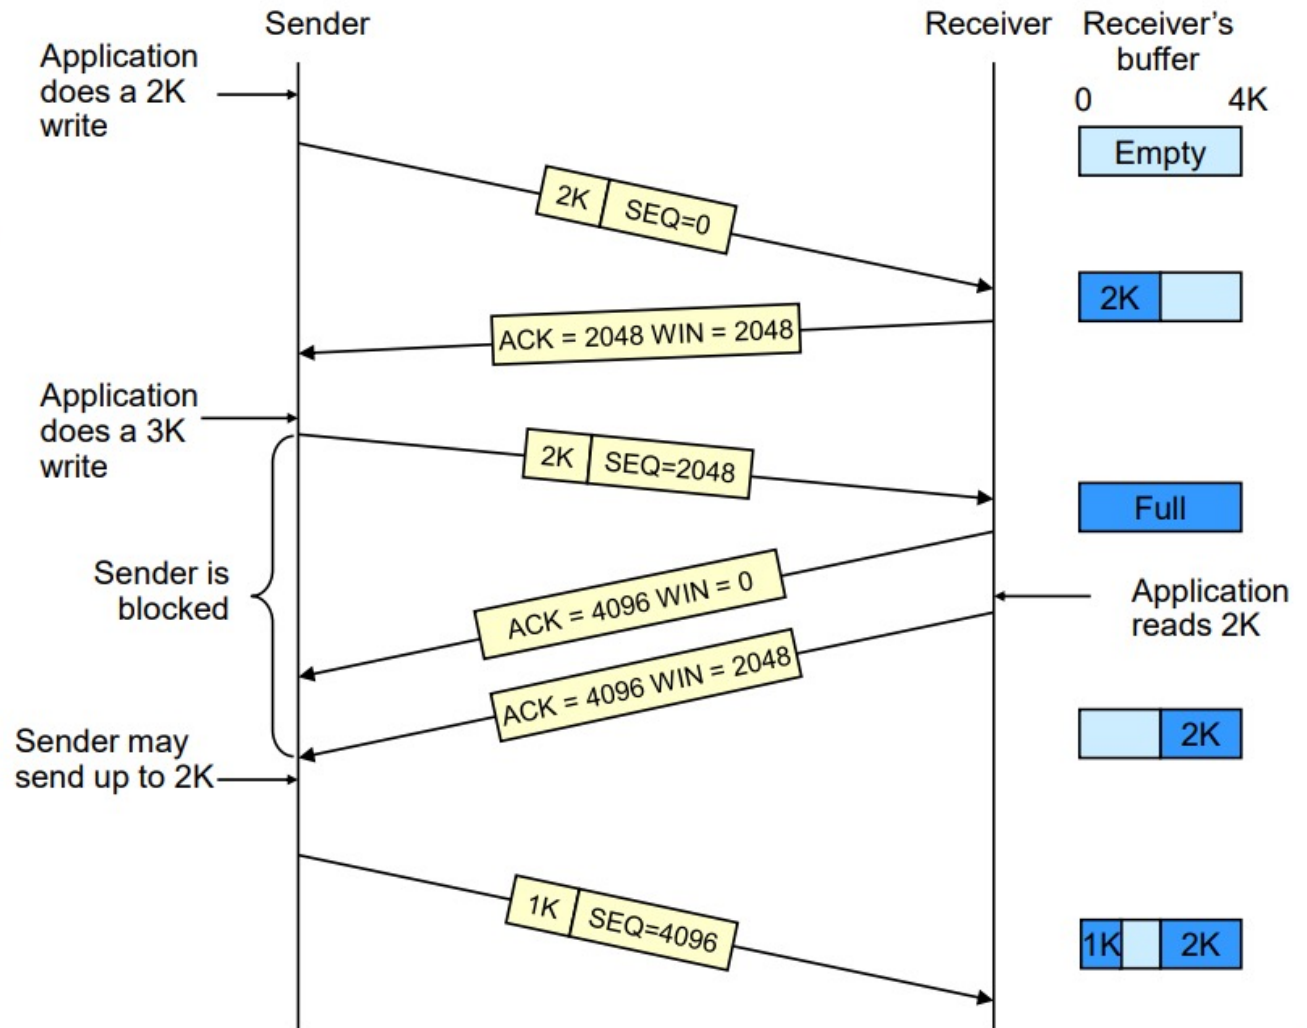
TCP Flow Control (cont'd)

Number of bytes in packet (N)

Sequence number of first data byte in packet (SEQ)

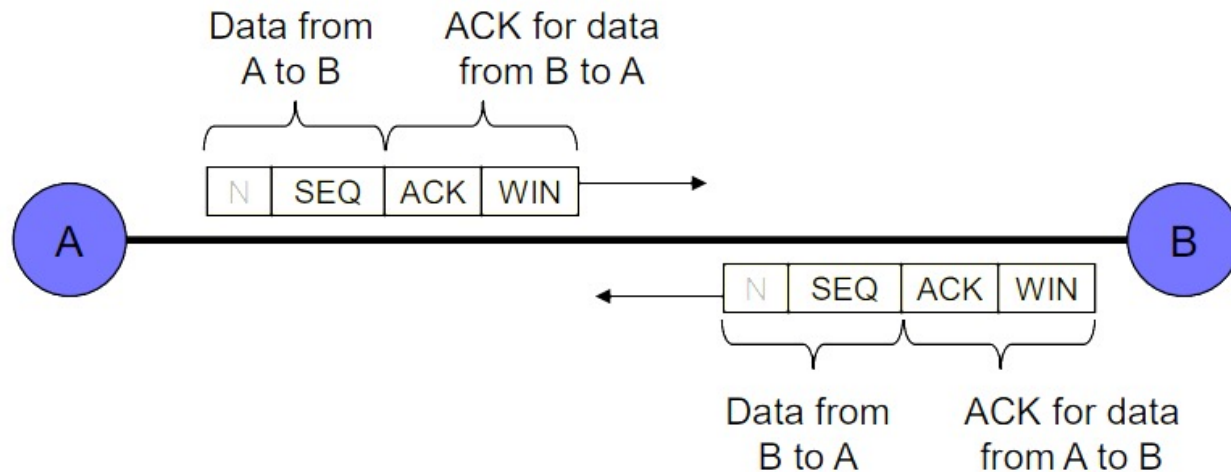
Sequence number of next expected byte (ACK)

Window size at the receiver (WIN)



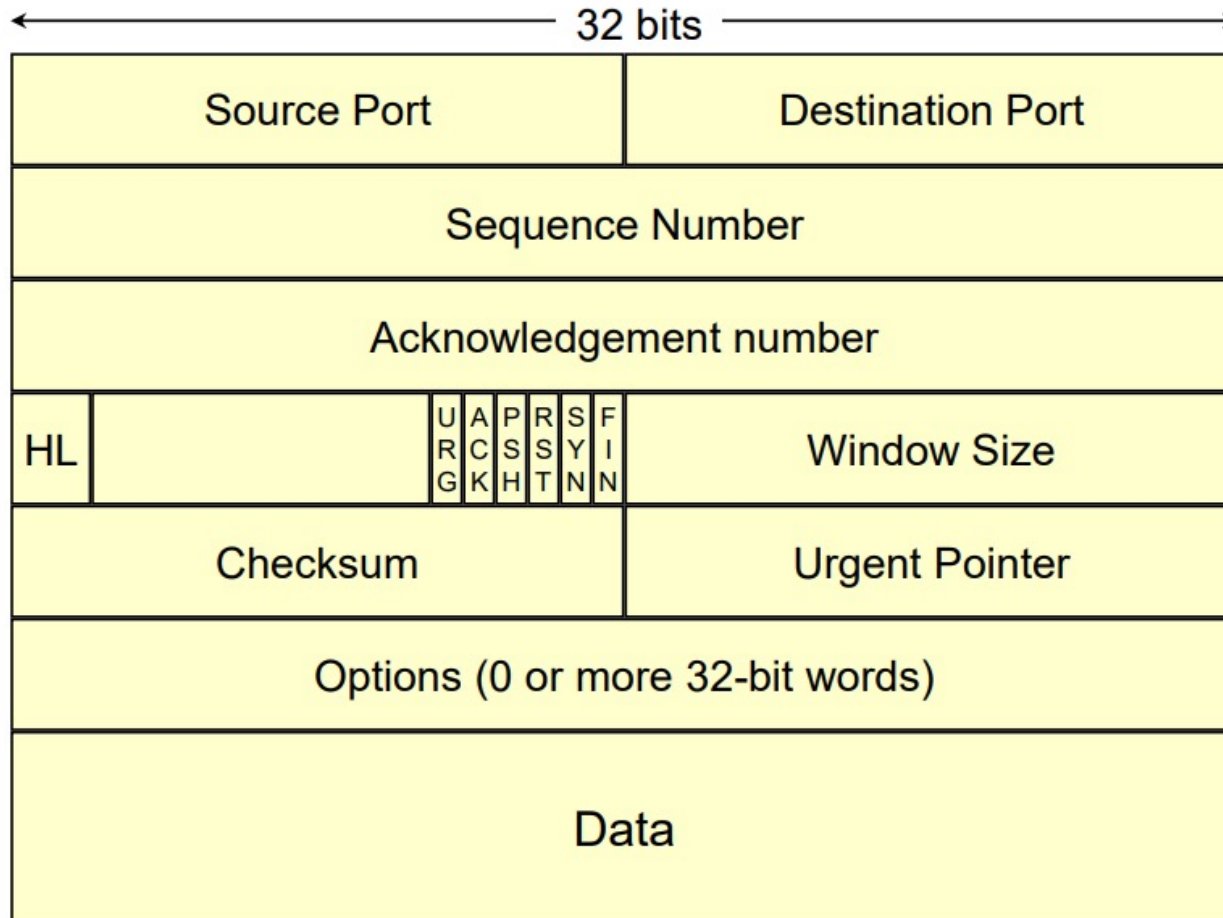
TCP Flow Control (cont'd)

Piggybacking: Allows more efficient bidirectional communication



4 TCP Frame Format

TCP Header Format



TCP Header Fields

- ❑ Source & Destination Ports
 - ❖ 16 bit port identifiers for each packet
- ❑ Sequence number
 - ❖ The packet's unique sequence ID
- ❑ Acknowledgement number
 - ❖ The sequence number of the next packet expected by the receiver

Source Port				Destination Port			
Sequence Number							
Acknowledgement number							
HL		U R G	A C K	P R S T	S Y N	F I N	Window Size
Checksum				Urgent Pointer			
Options (0 or more 32-bit words)							
Data							

TCP Header Fields (cont'd)

❑ Window size

- ❖ Specifies how many bytes may be sent after the first acknowledged byte

❑ Checksum

- ❖ Checksums the TCP header and IP address fields

❑ Urgent Pointer

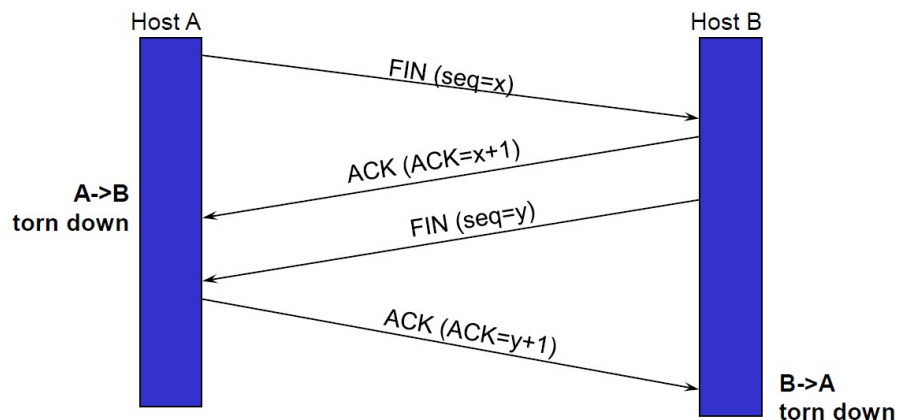
- ❖ Points to urgent data in the TCP data field

Source Port				Destination Port											
Sequence Number															
Acknowledgement number															
HL		URG		ACK		RST		SYN		FIN		Window Size			
Checksum						Urgent Pointer									
Options (0 or more 32-bit words)															
Data															

TCP Header Fields (cont'd)

□ Header bits

- ❖ URG = Urgent pointer field in use
- ❖ ACK = Indicates whether frame contains acknowledgement
- ❖ PSH = Data has been "pushed". It should be delivered to higher layers right away.
- ❖ RST = Indicates that the connection should be reset
- ❖ SYN = Used to establish connections
- ❖ FIN = Used to release a connection



Source Port				Destination Port													
Sequence Number																	
Acknowledgement number																	
HL		URG		ACK		PSH		RST		SYN		FIN		Window Size			
		GK		KH		TN		NN						Urgent Pointer			
Checksum								Options (0 or more 32-bit words)									
Data																	

5 Congestion Control

Principles of Congestion Control

Congestion:

- ❑ informally: “too many sources sending too much data too fast for *network* to handle”
- ❑ different from flow control!
- ❑ manifestations:
 - ❖ lost packets (buffer overflow at routers)
 - ❖ long delays (queueing in router buffers)

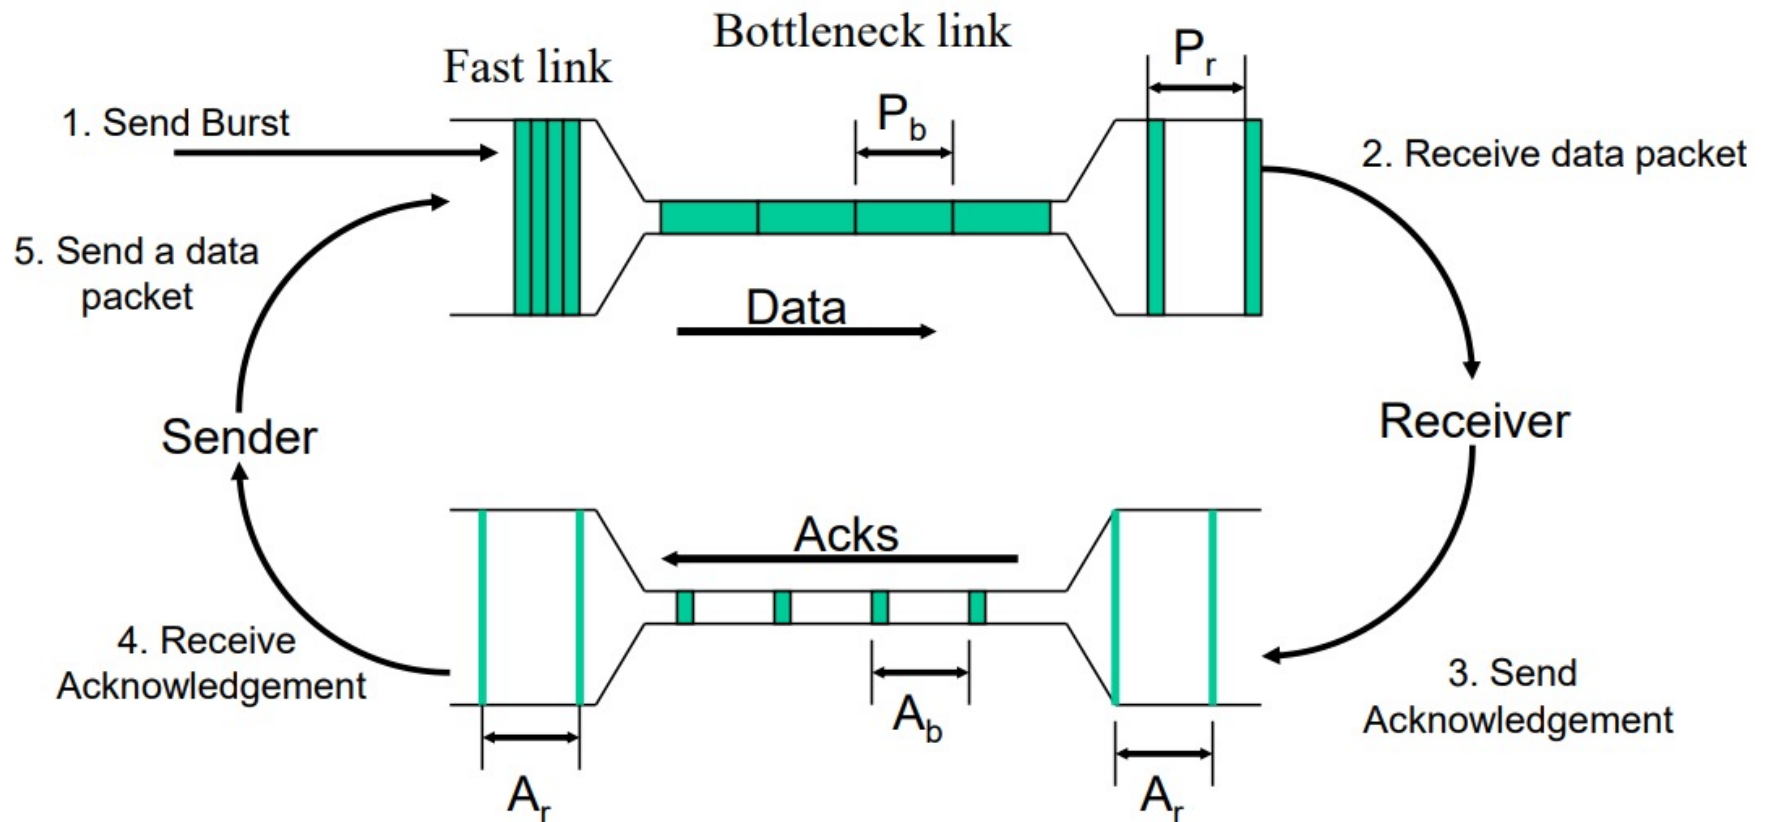
TCP Congestion Control

- ❑ Recall: Network layer is responsible for congestion control
- ❑ However, TCP/IP blurs the distinction
- ❑ In TCP/IP:
 - ❖ the network layer (IP) simply handles routing and packet forwarding
 - ❖ congestion control is done end-to-end by TCP

TCP Congestion Control

- Goal: fully (fairly) utilize the resource (bandwidth)
 - ❖ Don't over use - congestion
 - ❖ Don't under use - waste
- Goal: achieve self-clocking state
 - ❖ Even if don't know bandwidth of bottleneck
 - ❖ Bottleneck may change over time

Self-Clocking Model



Given: $P_b = P_r = A_r = A_b = A_r$ (in units of time)

Sending a packet on each ACK keeps the bottleneck link busy

TCP Congestion Window

- ❑ TCP introduces a second window, called the "congestion window"
- ❑ This window maintains TCP's best **estimate** of amount of outstanding data to allow in the network to achieve self-clocking
- ❑ Sending size = $\min(\text{congestion control window, flow control window})$

TCP Congestion Control

- Two phases to keep bottleneck busy (fully utilize the resource):
 - ❖ Increase the usage (window size) to keep probing the network
 - ❖ Decrease the usage when congestion is detected

TCP Slow Start


- When connection begins, $\text{CongWin} = 1 \text{ MSS}$
 - ❖ Example: $\text{MSS} = 500 \text{ bytes}$
- available bandwidth may be $\gg \text{MSS}/\text{RTT}$
 - ❖ desirable to quickly ramp up to respectable rate
 - ❖ Increase exponentially until first loss

MSS - "maximum segment size",
the maximum size a TCP packet
can be (including header)

TCP Slow Start (more)

- ❖ incrementing CongWin for every ACK received
- ❖ double CongWin every RTT

initial rate is slow but ramps up exponentially fast



TCP Slow Start (cont'd)


❑ Congestion detection

- ❖ Packet losses
- ❖ Sender side Timeout

❑ Timeout

- ❖ the congestion window is reduced to 1 MSS
- ❖ everything starts over

TCP Slow Start (cont'd)



TCP Linear Increase

- ❑ Don't push the network too fast
- ❑ Slow start (exponential)
 - > Threshold -> linear increase


TCP Linear Increase Algorithm

□ Algorithm:

- ❖ Start the threshold at 64K
- ❖ Slow start
- ❖ Once the threshold is passed
 - For each ack received, $cwnd = cwnd + (mss * mss) / cwnd$
 - 1 MSS for each congestion window of data transmitted
 - Timeout
 - reset the congestion window size to 1 MSS
 - Set threshold to $\max(2 * mss, 1/2 \text{ of MIN(sliding window, congestion window)})$

TCP Linear Increase Threshold Phase

Example: Maximum segment size = 1K
Assume thresh=32K



TCP Congestion Control

- ❑ Can we do better at detecting congestion than using timeout?
- ❑ Receiver send duplicate ack for out-of-order packets
 - ❖ Possible loss?

TCP Fast Retransmit

- ❑ Idea: When sender sees 3 duplicate ACKs, it assumes something went wrong
- ❑ The packet is immediately retransmitted instead of waiting for it to timeout

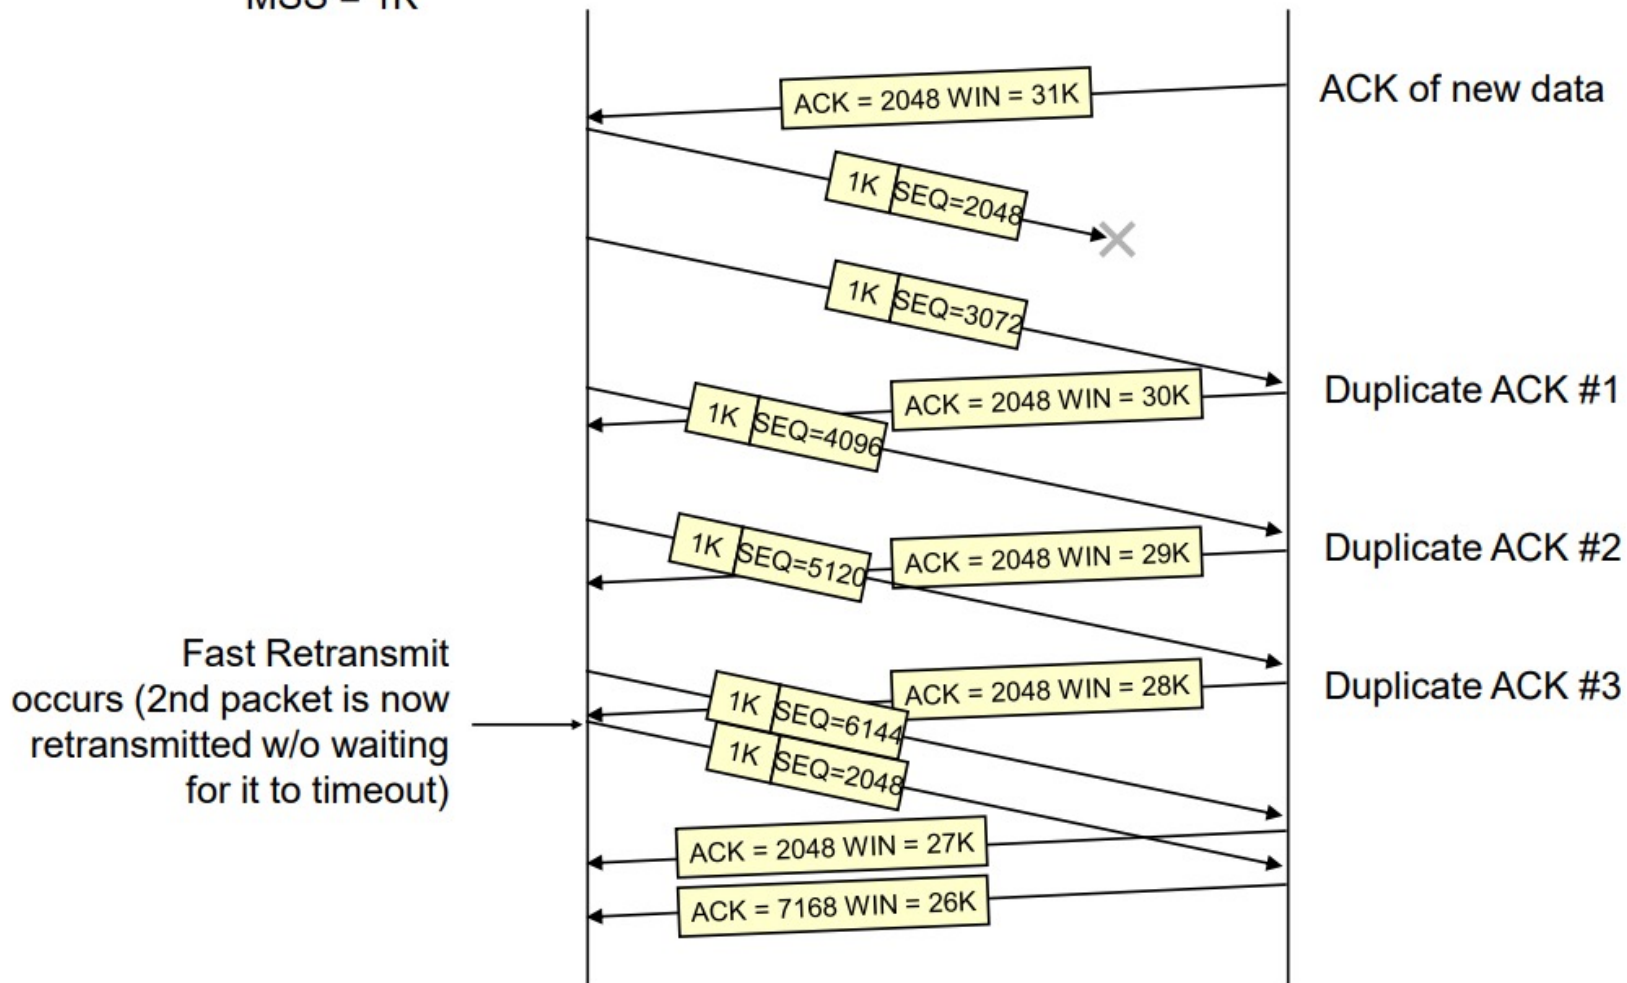
TCP Fast Retransmit

Example

MSS = 1K

Sender

Receiver



TCP Recap

❑ Timeout Computation

- ❖ Timeout is a function of 2 values
 - the weighted average of sampled RTTs
 - The sampled variance of each RTT

❑ Congestion control:

- ❖ Goal: Keep the self-clocking pipe full in spite of changing network conditions
- ❖ 3 key Variables:
 - Sliding window (Receiver flow control)
 - Congestion window (Sender flow control)
 - Threshold (Sender's slow start vs. linear mode line)

❑ Slow start

❑ Linear mode (Congestion Avoidance)

Algorithm Summary: TCP Congestion Control

- ❑ When CongWin is below Threshold, sender in **slow-start** phase, window grows exponentially.
- ❑ When CongWin is above Threshold, sender is in **congestion-avoidance** phase, window grows linearly.
- ❑ When a triple duplicate ACK occurs, Threshold set to $\max(\text{FlightSize}/2, 2 \cdot \text{mss})$ and CongWin set to $\text{Threshold} + 3 \cdot \text{mss}$. (**Fast retransmit, Fast recovery**)
- ❑ When timeout occurs, Threshold set to $\max(\text{FlightSize}/2, 2 \cdot \text{mss})$ and CongWin is set to 1 MSS.

FlightSize: The amount of data that has been sent but not yet acknowledged.