

# DNS

Why?

For any networked application, we need to know  
the IP address of a host given its name



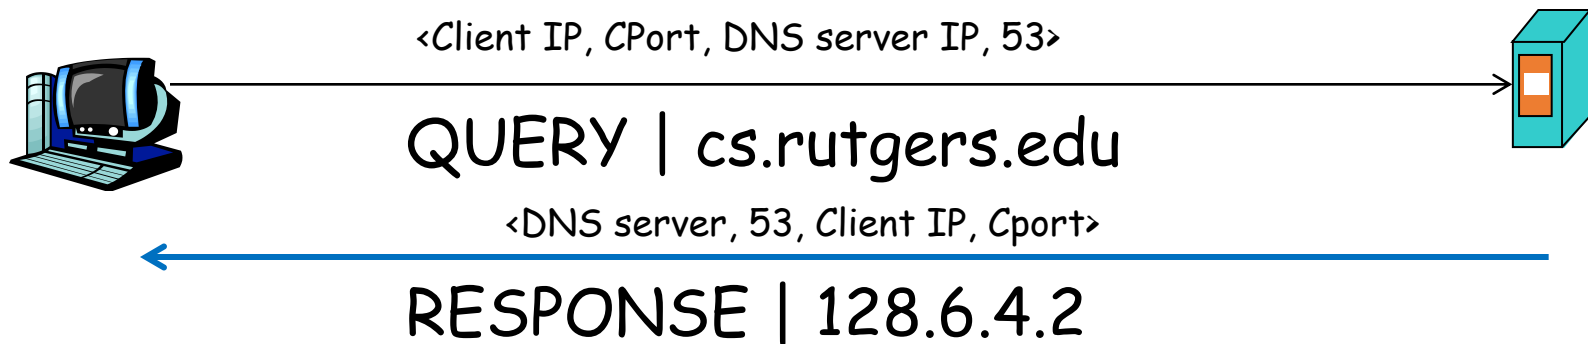
# Domain Name System (DNS)

- Problem statement:
  - Average brain can easily remember **7 digits** for a few names
  - On average, IP addresses have **12 digits**
  - We need an **easier** way to remember IP addresses
- Solution:
  - Use **names** to refer to hosts
  - Just as a contact or telephone **book** (white pages)
  - Add a **service** (called DNS) to map between host names and binary IP addresses
  - We call this ***Address Resolution***



# Simple DNS

DOMAIN NAME	IP ADDRESS
WWW.YAHOO.COM	98.138.253.109
cs.rutgers.edu	128.6.4.2
www.google.com	74.125.225.243
www.princeton.edu	128.112.132.86



- Simple but does not scale
- Every new host needs to be entered in this table
- Performance? Failure?

# DNS

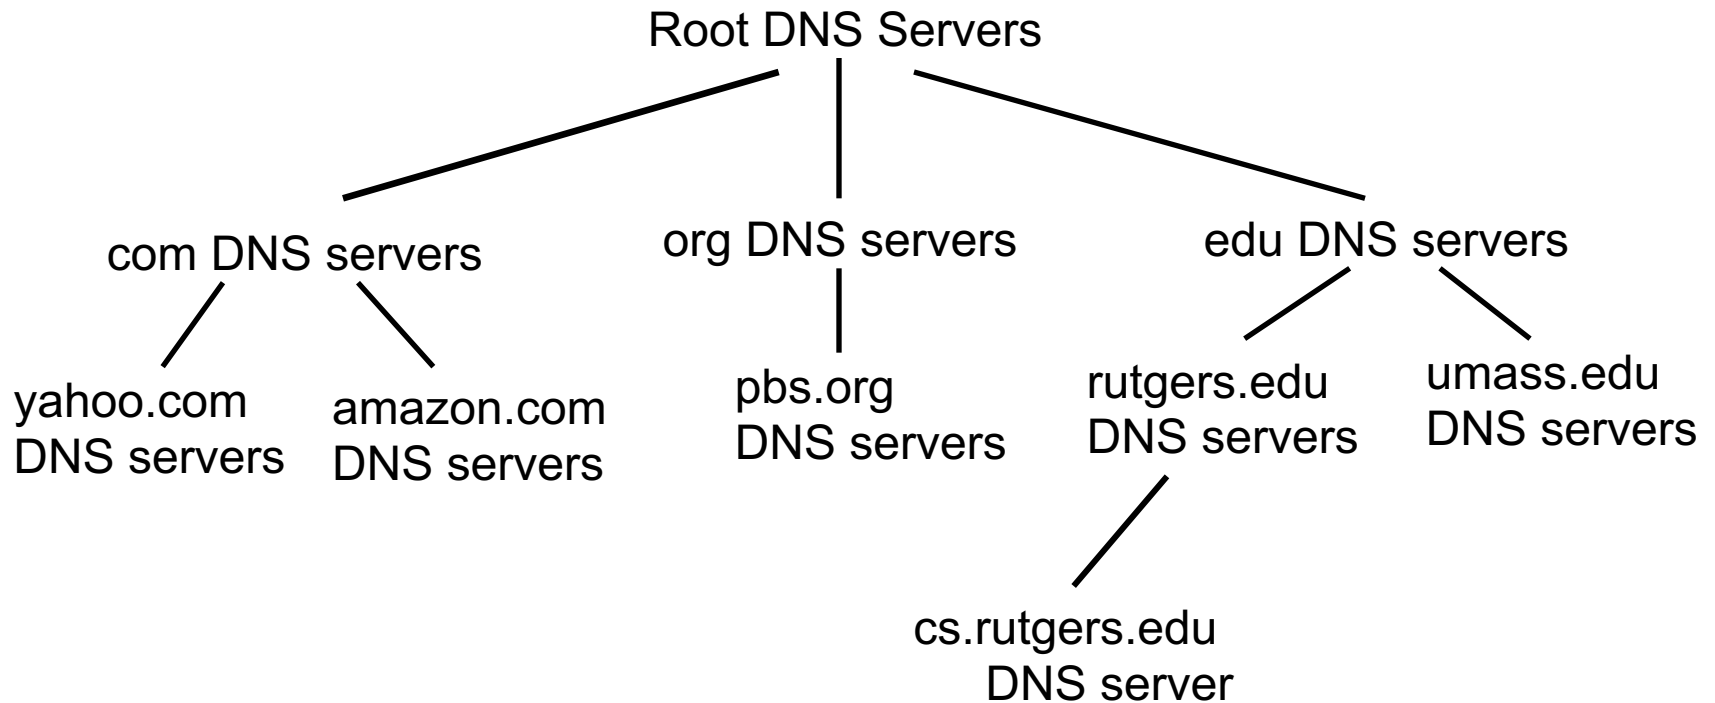
## Centralize DNS?

- single point of failure
- traffic volume
- Distant centralized database
- maintenance

doesn't *scale!*



# Distributed, Hierarchical Database



RFC 1034



# DNS Protocol

- Client and Server (CS Model)
- Client connects to **Port 53**
- DNS server address should be **known**
  - Either manually configured or automatically
- Two types of messages
  - Queries
  - Responses
- Type of Query methods
  - **Standard query**
    - Request IP address
  - **Updates**
    - Provide a binding of IP address to domain name
- Each type has a **common message format** that follows the header




# DNS Protocol

- When client wants to know an IP address for a host name
  - Client **sends** a DNS query to the primary name server in its zone
  - If name server **contains the mapping**, it returns the IP address to the client
  - Otherwise, the name server **forwards the request** to the root name server
  - The request **works its way down the tree** toward the host until it **reaches a name server** with the correct mapping



# Example

- Host at **cis.poly.edu** wants IP address for **gaia.cs.umass.edu**






# Query type

## iterated query:


- contacted server **replies** with name of server to contact
- “I don’t know this name, but ask this server”



# Query type

## recursive query:

- puts burden of name resolution on contacted name server



# DNS: caching and updating records

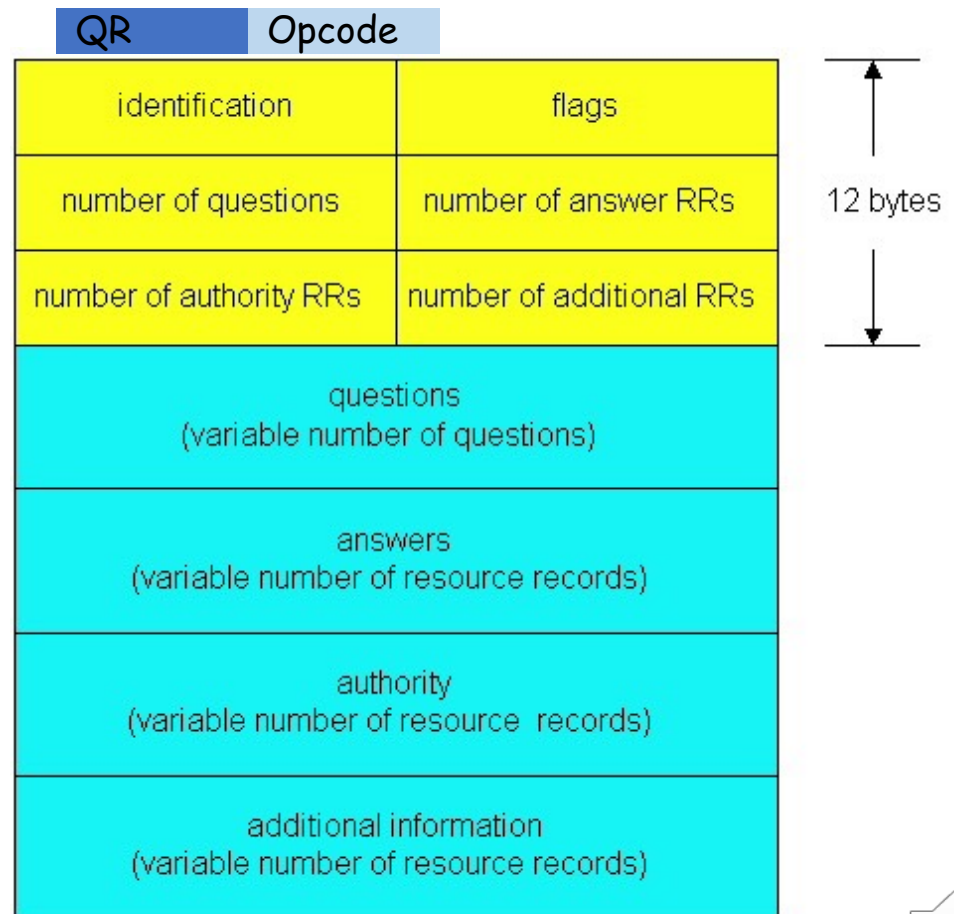
- once (any) name server learns mapping, it *caches* mapping
  - cache entries **timeout** (disappear) after some time
- TLD (Top Level Domain) servers typically **cached** in local name servers
  - Thus root name servers not often **visited**

# DNS protocol, messages

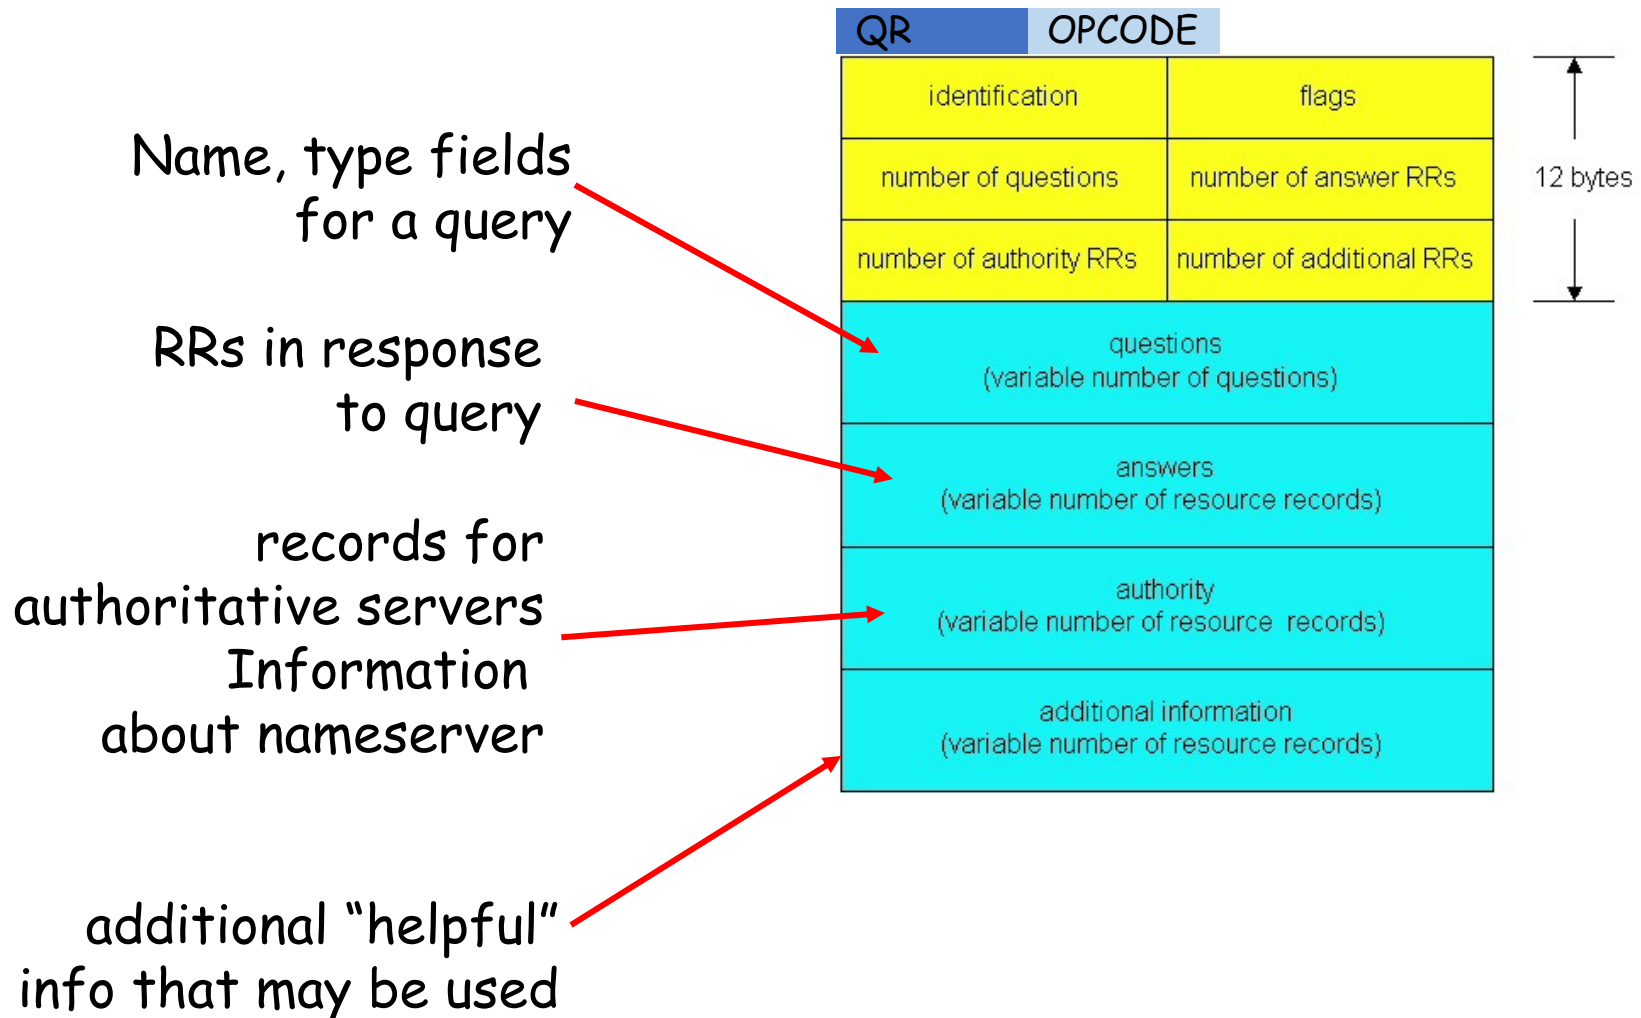
DNS protocol : *query* and *reply* messages, both with same *message format*

## msg header

- ❑ QR = 0 for Q, 1 for response
- ❑ Opcode= 0 standard
- ❑ identification: 16 bit # for query, reply to query uses same #
- ❑ flags:
  - ❖ Authoritative answer
  - ❖ recursion desired
  - ❖ recursion available
  - ❖ reply is authoritative



# DNS protocol, messages



# DNS records

DNS: distributed db storing resource records (RR)

RR format: (name, **type**, class, ttl, addr)

## □ Type=A

- ❖ **name** is hostname
- ❖ **value** is IP address

## • Type=NS

- **name** is domain (e.g. foo.com)
- **value** is hostname of **authoritative name server** for this domain

## □ Type=AAAA

- ❖ **name** is hostname
- ❖ **value** is IPv6 address

# DNS Record example

RRs in response  
to query



NAME	Design.cs.rutgers.edu
TYPE	A
CLASS	IN
TTL	1 day(86400)
ADDRESS	192.26.92.30

records for  
authoritative  
servers

Information about  
nameserver



NAME	Cs.rutgers.edu
TYPE	NS
CLASS	IN
TTL	1 day(86400)
NSD NAME	Ns-lcsr.rutgers.edu

# Bootstrapping DNS

- How does a host **contact** the name server if **all it has is the name** and **no IP address**?
- IP address of at least 1 nameserver must be given in advance
  - or with another protocol (DHCP, bootp)
- File `/etc/resolv.conf` in unix
- Start -> settings-> control panel-> network -> TCP/IP -> properties **in windows**



# Themes

- Request/response nature of these protocols
- How Messages are structured
  - HTTP, SMTP, FTP - simple ASCII protocols
- Caching
- Name Lookup
  - Hierarchy structure



# HTTP

Hypertext Transfer Protocol



# Web and HTTP

## First some jargon

- **Web page** consists of **objects**
- Object can be HTML file, JPEG image, Java applet, audio file,...
- Web page consists of **base HTML-file** which includes several referenced objects
- Each object is addressable by a **URL ( Uniform Resource Locator)**
- Example URL:


`www.cs.rutgers.edu/undergraduate/pic.gif`

host name                      path name

# HTTP overview

## HTTP: hypertext transfer protocol

- client/server model
  - *client*: browser that requests, receives, “displays” Web objects
  - *server*: Web server sends objects in response to requests
- HTTP 1.0: RFC 1945
- HTTP 1.1: RFC 2068



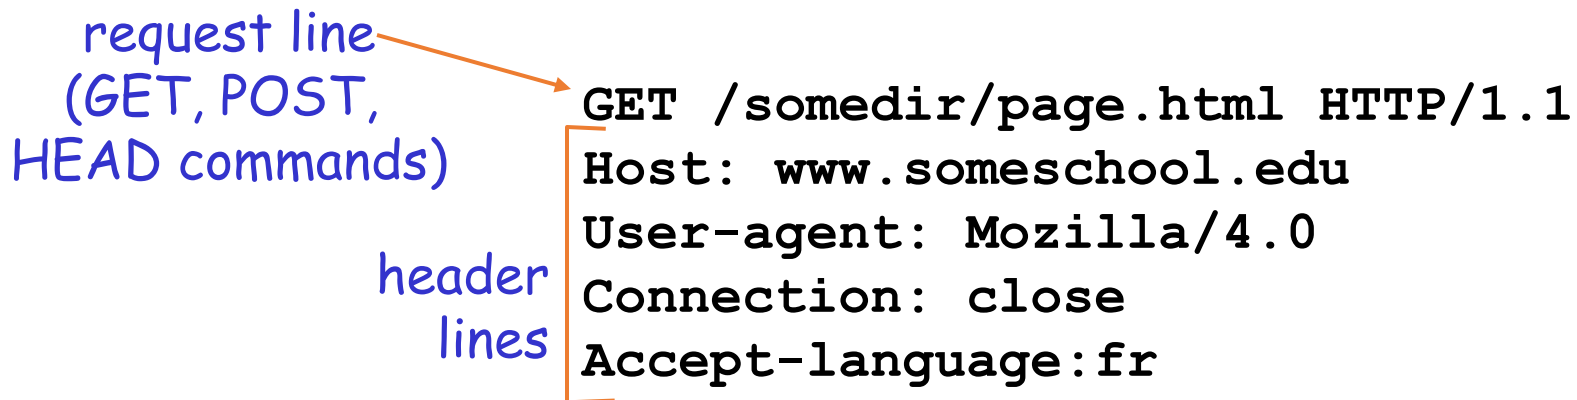
# HTTP messages: request message

- HTTP request message:
  - ASCII (human-readable format)

request line  
(GET, POST,  
HEAD commands)


header  
lines

```
GET /somedir/page.html HTTP/1.1
Host: www.someschool.edu
User-agent: Mozilla/4.0
Connection: close
Accept-language: fr
```

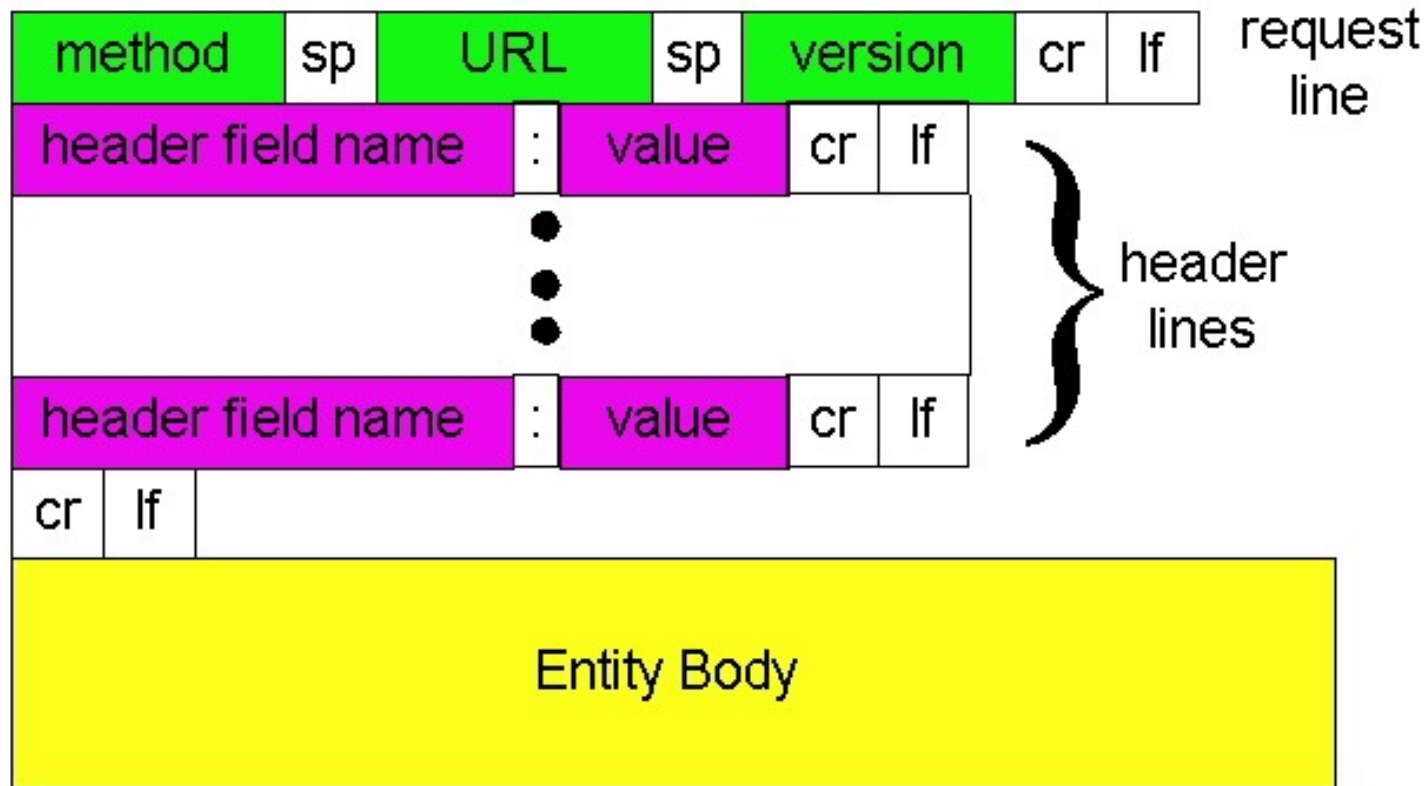


# Client server connection

DNS



# HTTP request message: general format



# Method types

- GET

- ❖ Get the file **specified** in the path URL field in entity body

- POST

- accept the entity enclosed in the entity body as **a new subordinate** of the resource identified by the URL field

- PUT

- uploads file in **entity body** to path specified in URL field


- DELETE

- deletes file specified in the URL field

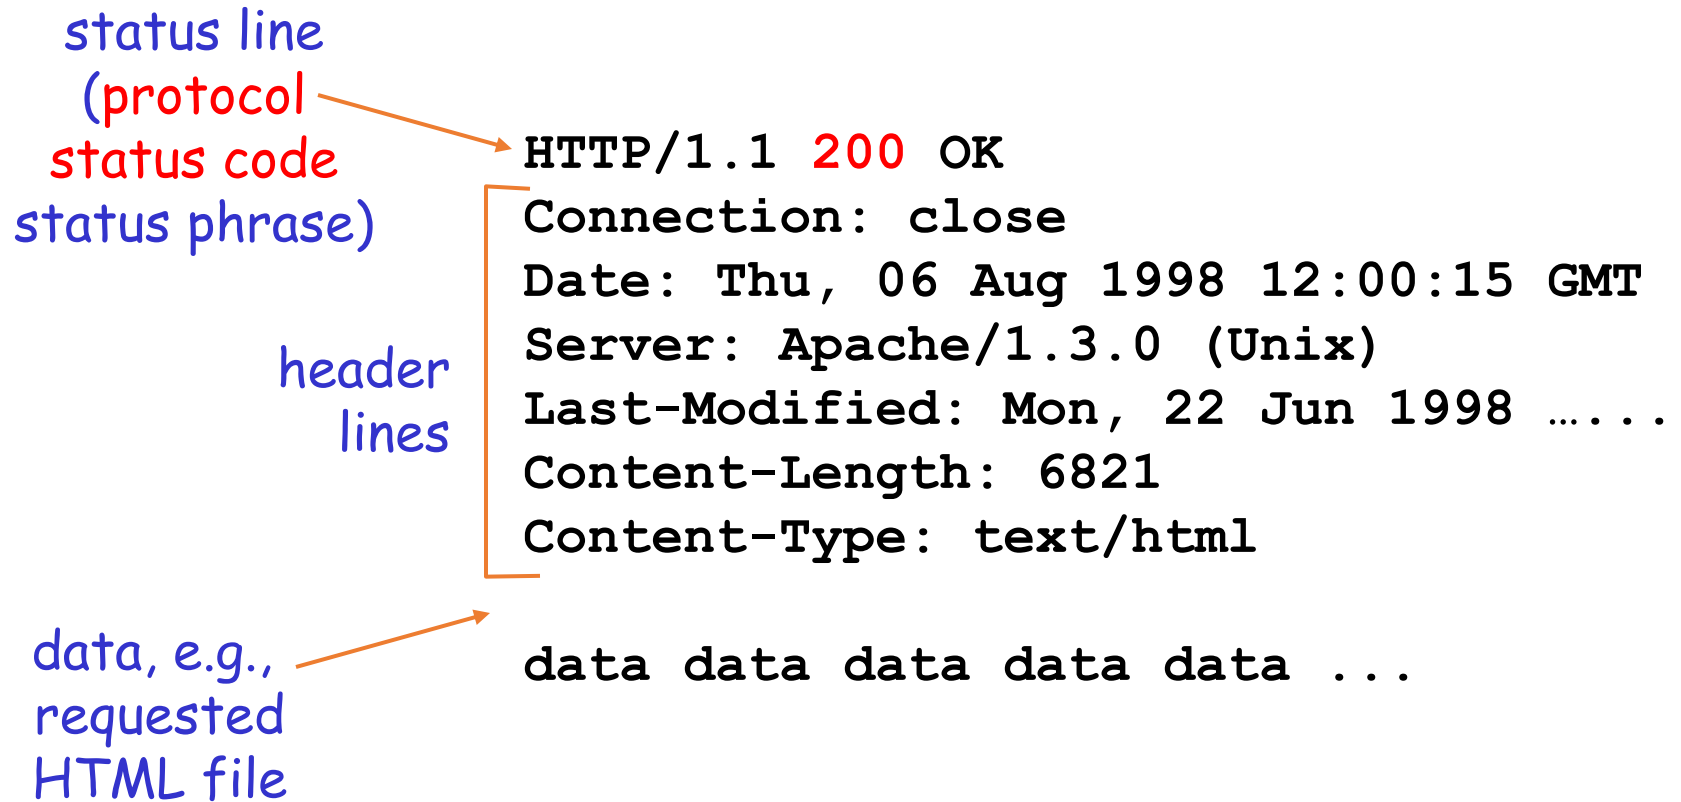


# http response message: general format

Unlike http request, No method name



# HTTP message: response message



# HTTP response status codes

In first line in server->client response message.

A few sample codes:

## 200 OK

- request succeeded, requested object later in this message

## 301 Moved Permanently

- requested object **moved**, new location **specified** later in this message (Location:)

## 400 Bad Request

- request message **not understood** by server

## 404 Not Found

- requested document **not found** on this server

## 505 HTTP Version Not Supported



# Additional about HTTP

- Persistent vs. Nonpersistent HTTP connections
- Cookies (User-server state)
- Web caches

# HTTP connections

## Nonpersistent HTTP

- At most **one object** is sent over a **single** TCP connection.
- HTTP/1.0 uses nonpersistent HTTP

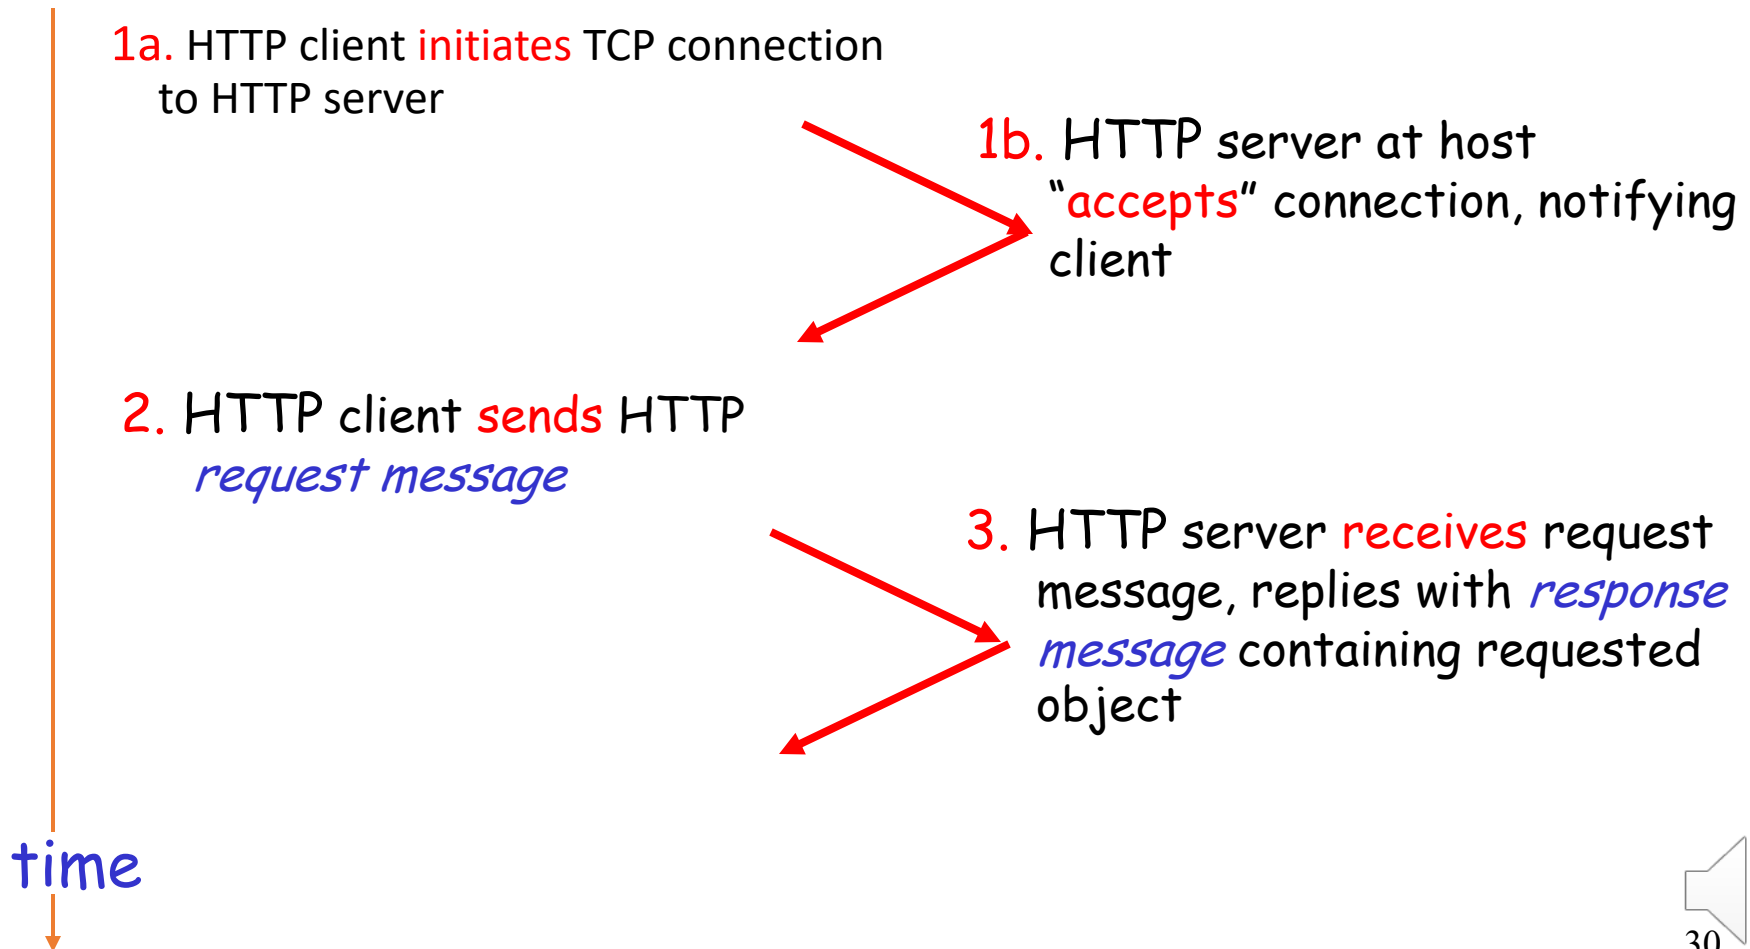
## Persistent HTTP

- **Multiple objects** can be sent over a **single** TCP connection between client and server.
- HTTP/1.1 uses persistent connections in default mode

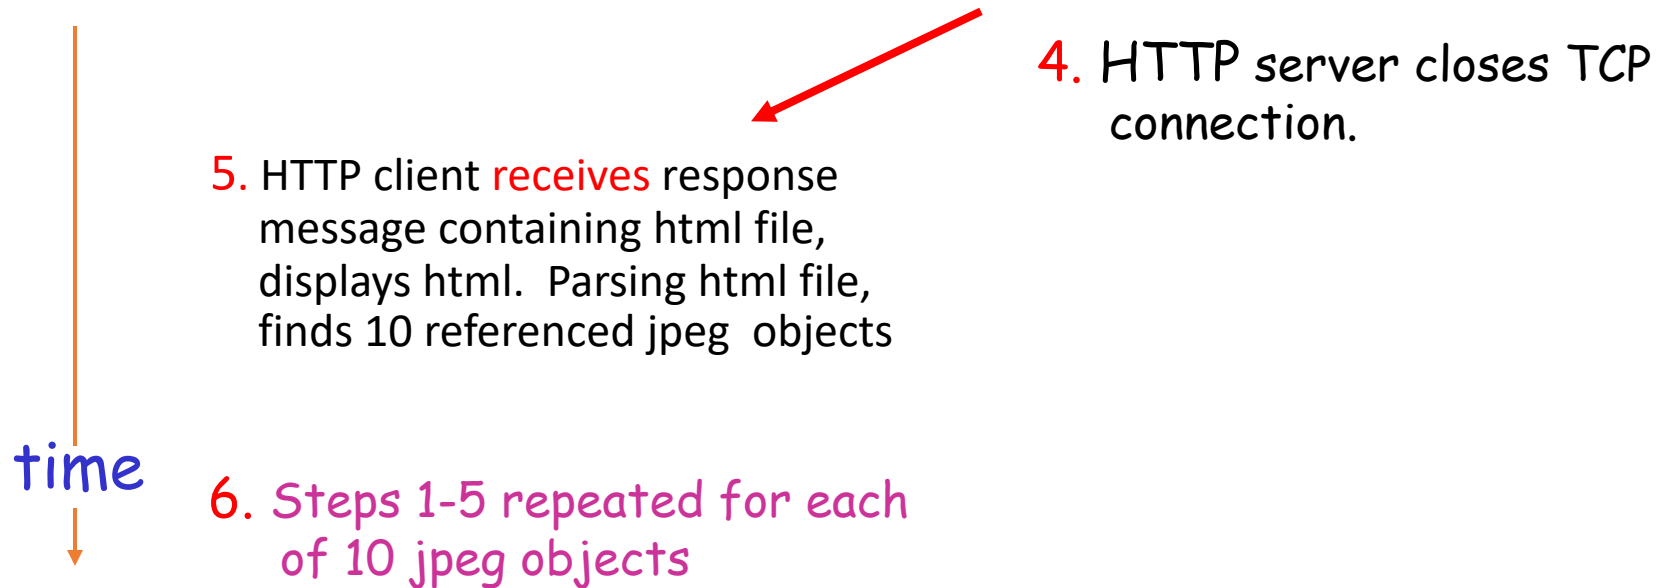
TCP is a kind of communication service provided by the transport layer. It requires the connection to be set up before data communication.

# Nonpersistent HTTP

Suppose user enters URL `www.someSchool.edu/someDepartment/home.index` (contains text, references to 10 jpeg images)



# Nonpersistent HTTP (cont.)




# HTTP: Response time

**Definition of RTT (Round-Trip Time):** time to send a small packet to travel from client to server and back.

## Response time:

- one RTT to **initiate** TCP connection
- one RTT for HTTP request and first few bytes of HTTP response to return
- file transmission time

**total =  $2RTT + \text{transmit time}$**





# Persistent vs. Nonpersistent

## Nonpersistent HTTP issues:

- requires 2 RTTs per object
  - TCP Connection and HTTP Request
- Browsers can open parallel TCP connections to fetch referenced objects

## Persistent HTTP

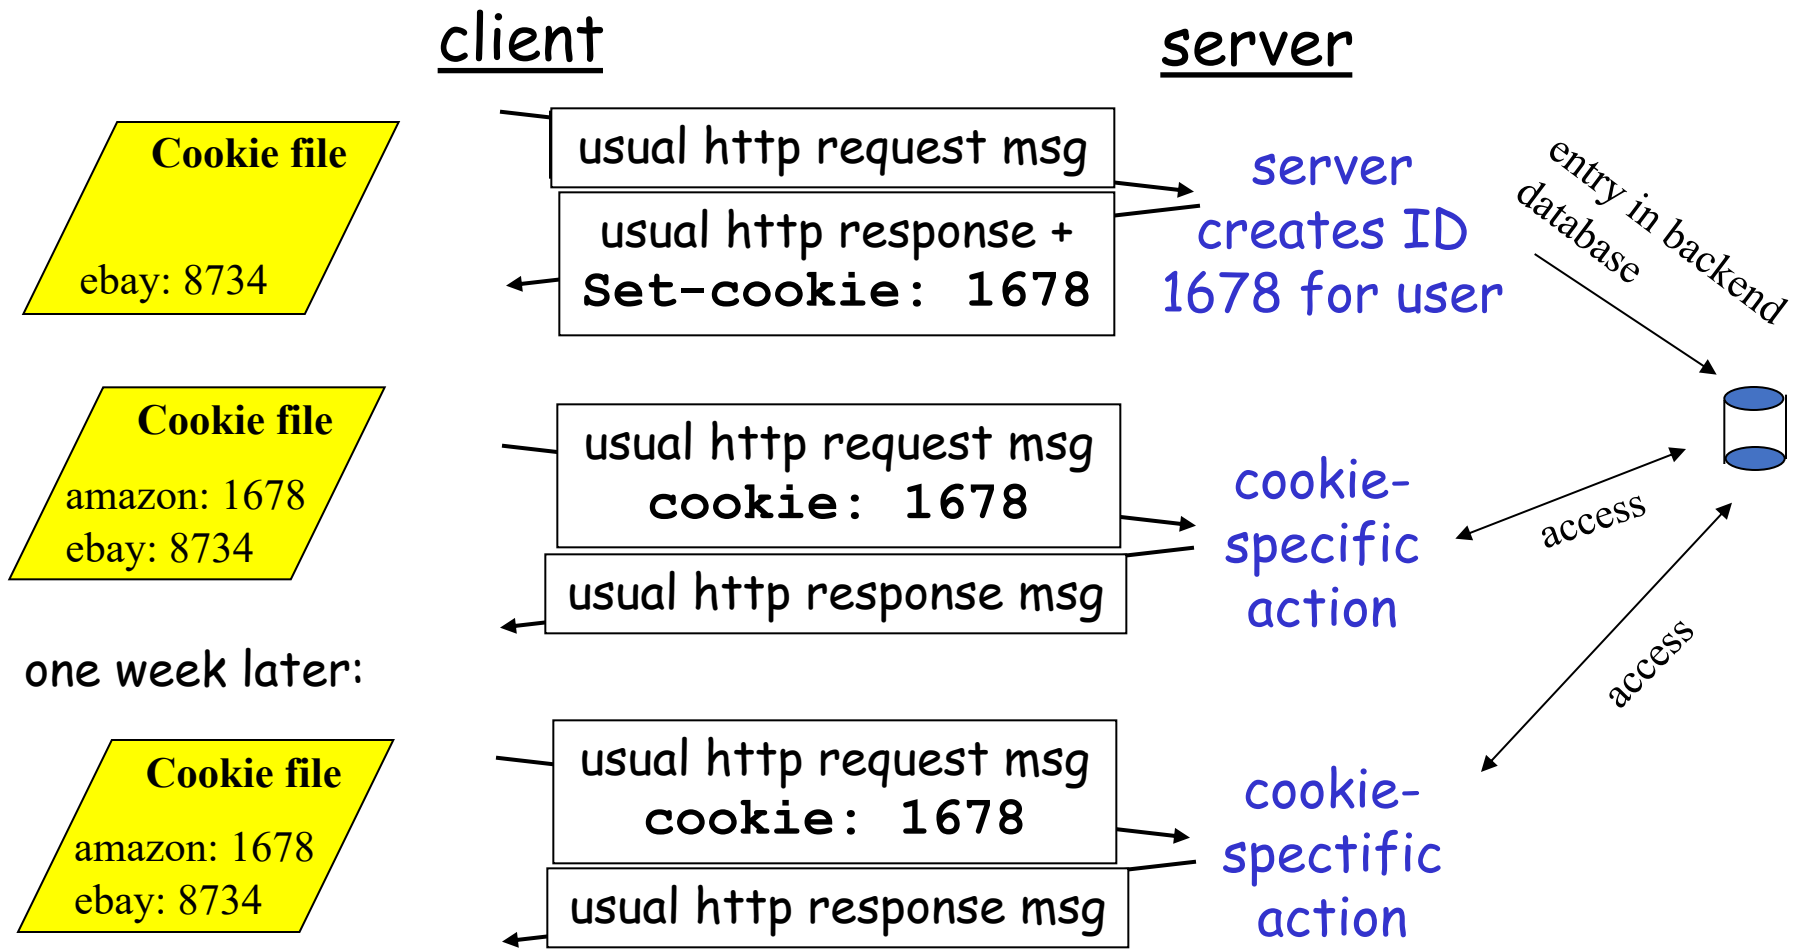
- server leaves TCP connection open after sending response
- subsequent HTTP messages between same client/server sent over open connection

# HTTP: user-server state

## HTTP is “stateless”

- server maintains no information about past client requests
- What **state** can bring:
  - authorization
  - shopping carts
  - recommendations
  - user session state

# Cookies: keeping “state”



# Cookies (continued)

## Four components:

- 1) cookie header line of HTTP *response* message
- 2) cookie header line in HTTP *request* message
- 3) cookie file kept on user's host, **managed** by user's browser
- 4) back-end database at Web site

# Cookies (continued)

— aside —  
Cookies and privacy:

- ❑ cookies permit sites to learn a lot about you

